

---

# MAC0352 - Redes de Computadores e Sistemas Distribuídos

Daniel Macêdo Batista

IME - USP, 13 de Maio de 2021

HTTP (Cookies)

Proxy

**HTTP (Cookies)**  
**Proxy**

▷ HTTP (Cookies)

Proxy

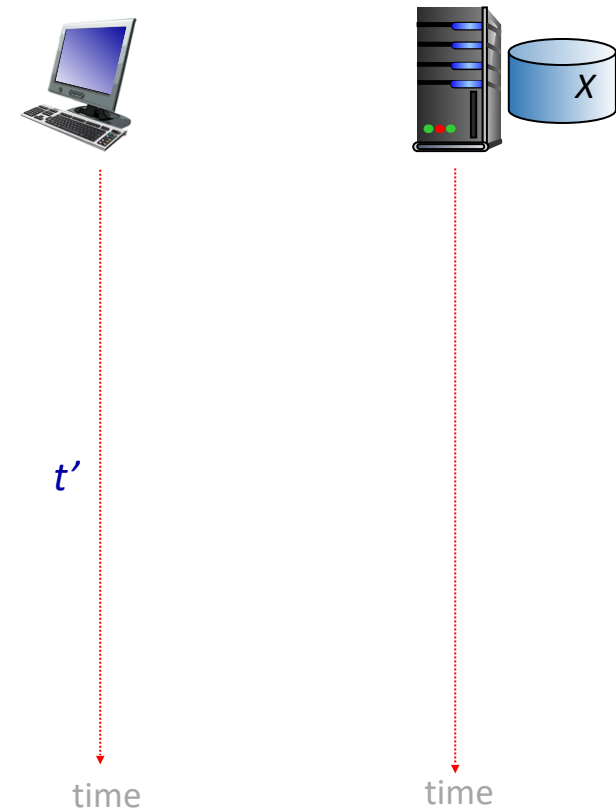
# HTTP (Cookies)

# Maintaining user/server state: cookies

Recall: HTTP GET/response  
interaction is *stateless*

- no notion of multi-step exchanges of HTTP messages to complete a Web “transaction”

a *stateful protocol*: client makes two changes to X, or none at all

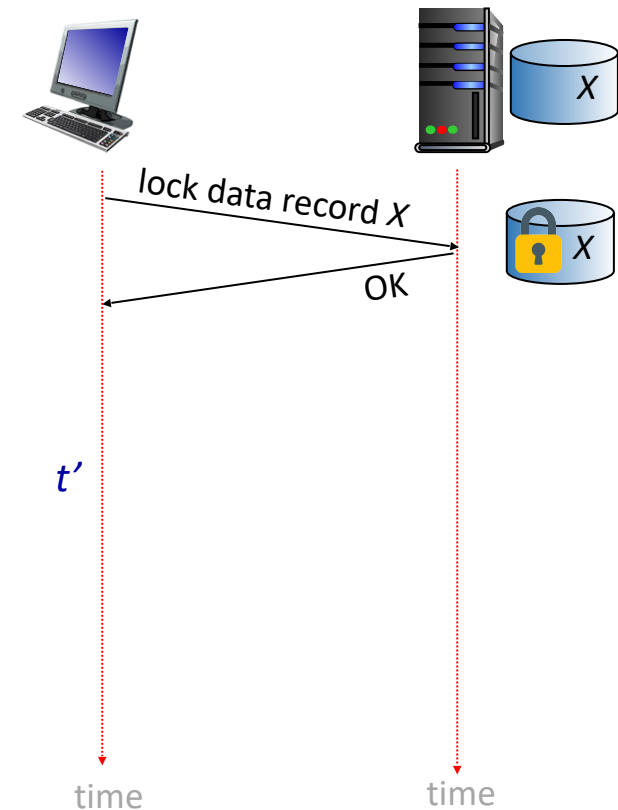


# Maintaining user/server state: cookies

Recall: HTTP GET/response  
interaction is *stateless*

- no notion of multi-step exchanges of HTTP messages to complete a Web “transaction”

a *stateful protocol*: client makes two changes to X, or none at all

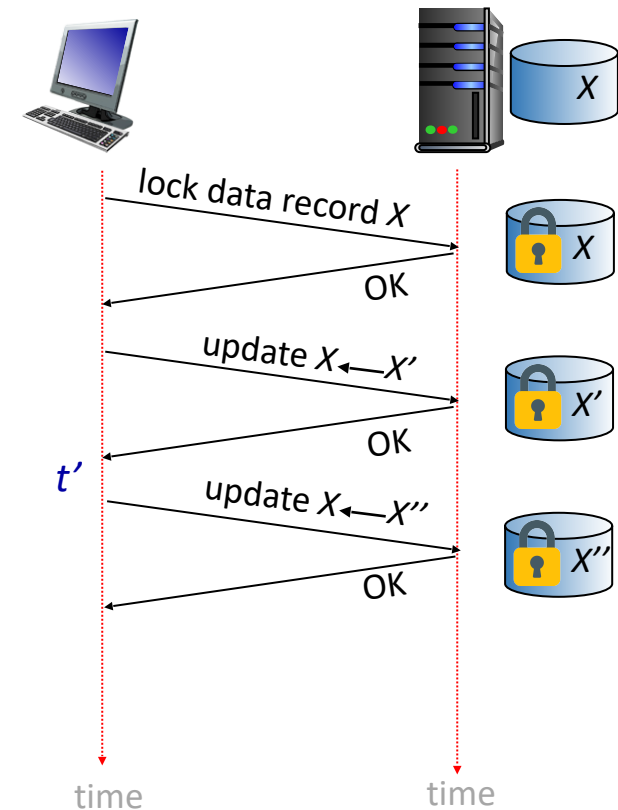


# Maintaining user/server state: cookies

Recall: HTTP GET/response  
interaction is *stateless*

- no notion of multi-step exchanges of HTTP messages to complete a Web “transaction”

a *stateful protocol*: client makes two changes to  $X$ , or none at all

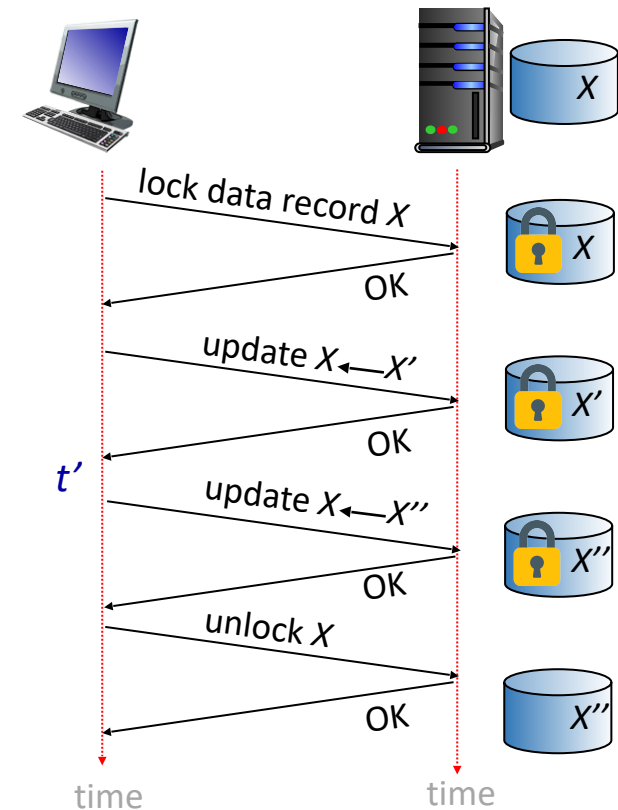


# Maintaining user/server state: cookies

Recall: HTTP GET/response  
interaction is *stateless*

- no notion of multi-step exchanges of HTTP messages to complete a Web “transaction”

a *stateful protocol*: client makes two changes to  $X$ , or none at all

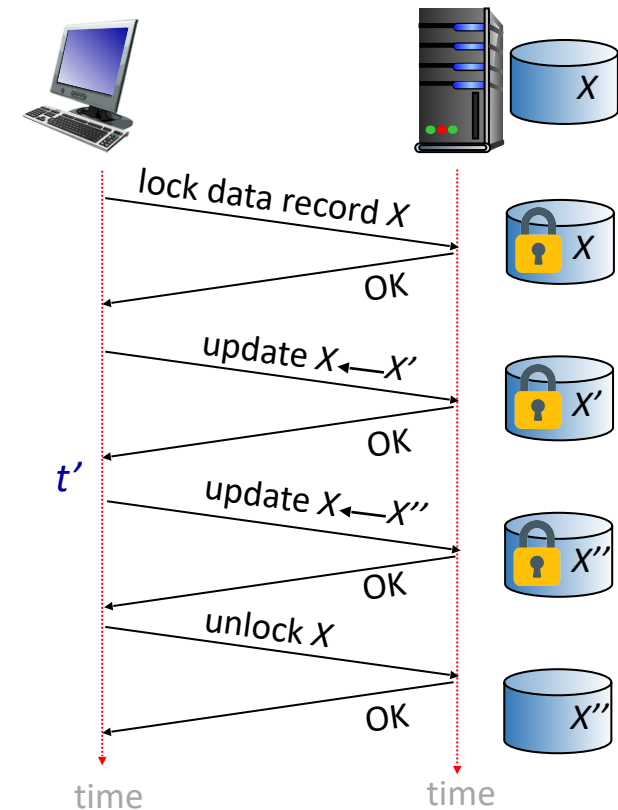


# Maintaining user/server state: cookies

Recall: HTTP GET/response  
interaction is *stateless*

- no notion of multi-step exchanges of HTTP messages to complete a Web “transaction”

a *stateful protocol*: client makes two changes to  $X$ , or none at all



*Q:* what happens if network connection or client crashes at  $t'$ ?

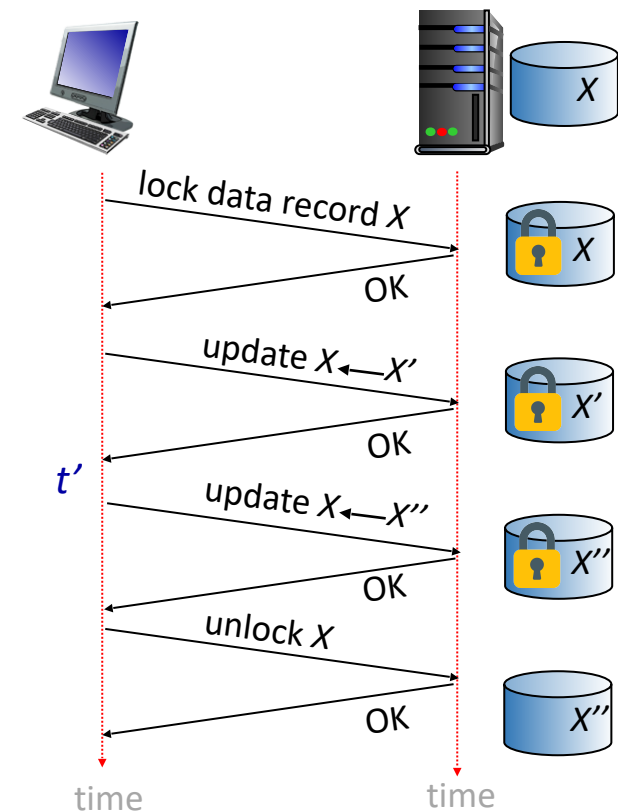


# Maintaining user/server state: cookies

Recall: HTTP GET/response interaction is *stateless*

- no notion of multi-step exchanges of HTTP messages to complete a Web “transaction”
  - no need for client/server to track “state” of multi-step exchange

a *stateful protocol*: client makes two changes to  $X$ , or none at all



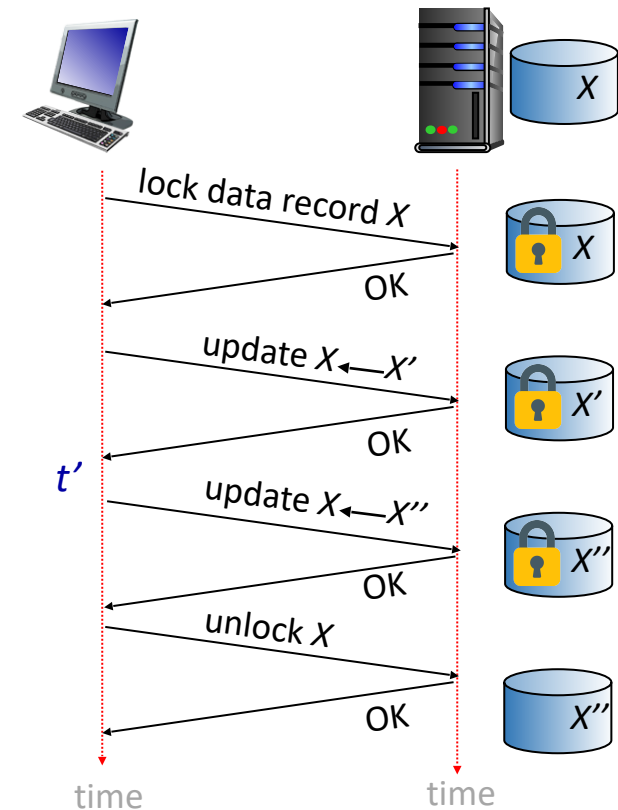
*Q:* what happens if network connection or client crashes at  $t'$ ?

# Maintaining user/server state: cookies

Recall: HTTP GET/response interaction is *stateless*

- no notion of multi-step exchanges of HTTP messages to complete a Web “transaction”
  - no need for client/server to track “state” of multi-step exchange
  - all HTTP requests are independent of each other

a *stateful protocol*: client makes two changes to  $X$ , or none at all



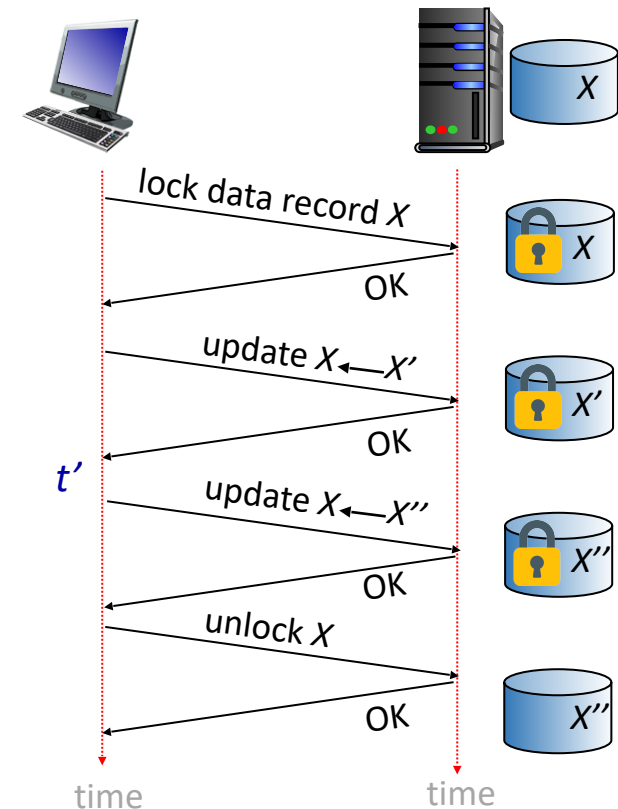
*Q:* what happens if network connection or client crashes at  $t'$ ?

# Maintaining user/server state: cookies

Recall: HTTP GET/response interaction is *stateless*

- no notion of multi-step exchanges of HTTP messages to complete a Web “transaction”
  - no need for client/server to track “state” of multi-step exchange
  - all HTTP requests are independent of each other
  - no need for client/server to “recover” from a partially-completed-but-never-completely-completed transaction

a *stateful protocol*: client makes two changes to  $X$ , or none at all



*Q:* what happens if network connection or client crashes at  $t'$ ?

# Maintaining user/server state: cookies

Web sites and client browser use *cookies* to maintain some state between transactions

*four components:*

# Maintaining user/server state: cookies

Web sites and client browser use *cookies* to maintain some state between transactions

*four components:*

- 1) cookie header line of HTTP *response* message

# Maintaining user/server state: cookies

Web sites and client browser use *cookies* to maintain some state between transactions

*four components:*

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in next HTTP *request* message

# Maintaining user/server state: cookies

Web sites and client browser use *cookies* to maintain some state between transactions

*four components:*

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in next HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser

# Maintaining user/server state: cookies

Web sites and client browser use *cookies* to maintain some state between transactions

*four components:*

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in next HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site



# Maintaining user/server state: cookies

Web sites and client browser use *cookies* to maintain some state between transactions

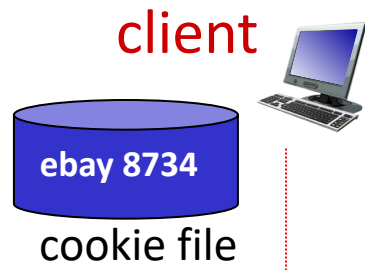
## *four components:*

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in next HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

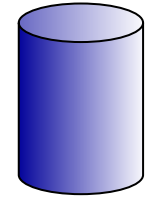
## Example:

- Susan uses browser on laptop, visits specific e-commerce site for first time
- when initial HTTP requests arrives at site, site creates:
  - unique ID (aka “cookie”)
  - entry in backend database for ID
- subsequent HTTP requests from Susan to this site will contain cookie ID value, allowing site to “identify” Susan

# Maintaining user/server state: cookies



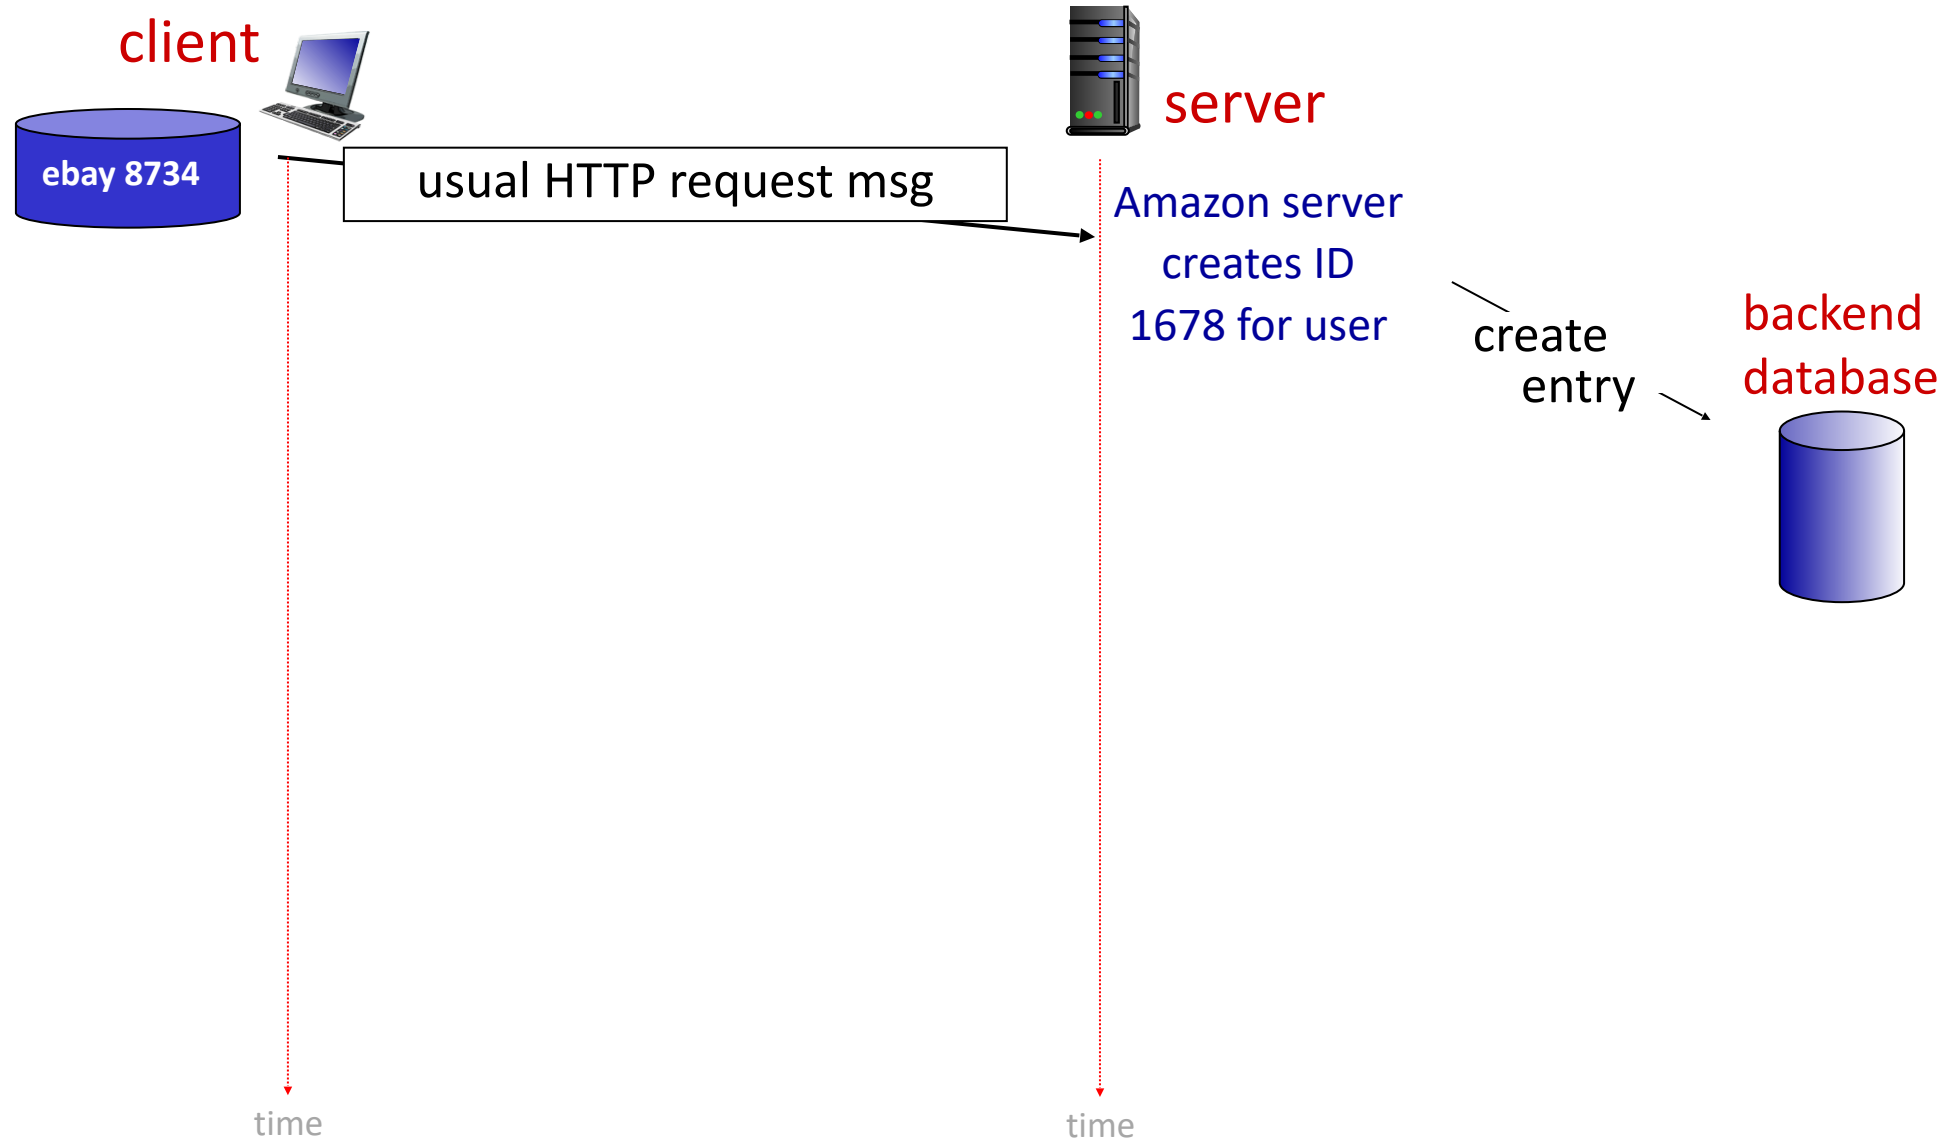
backend  
database



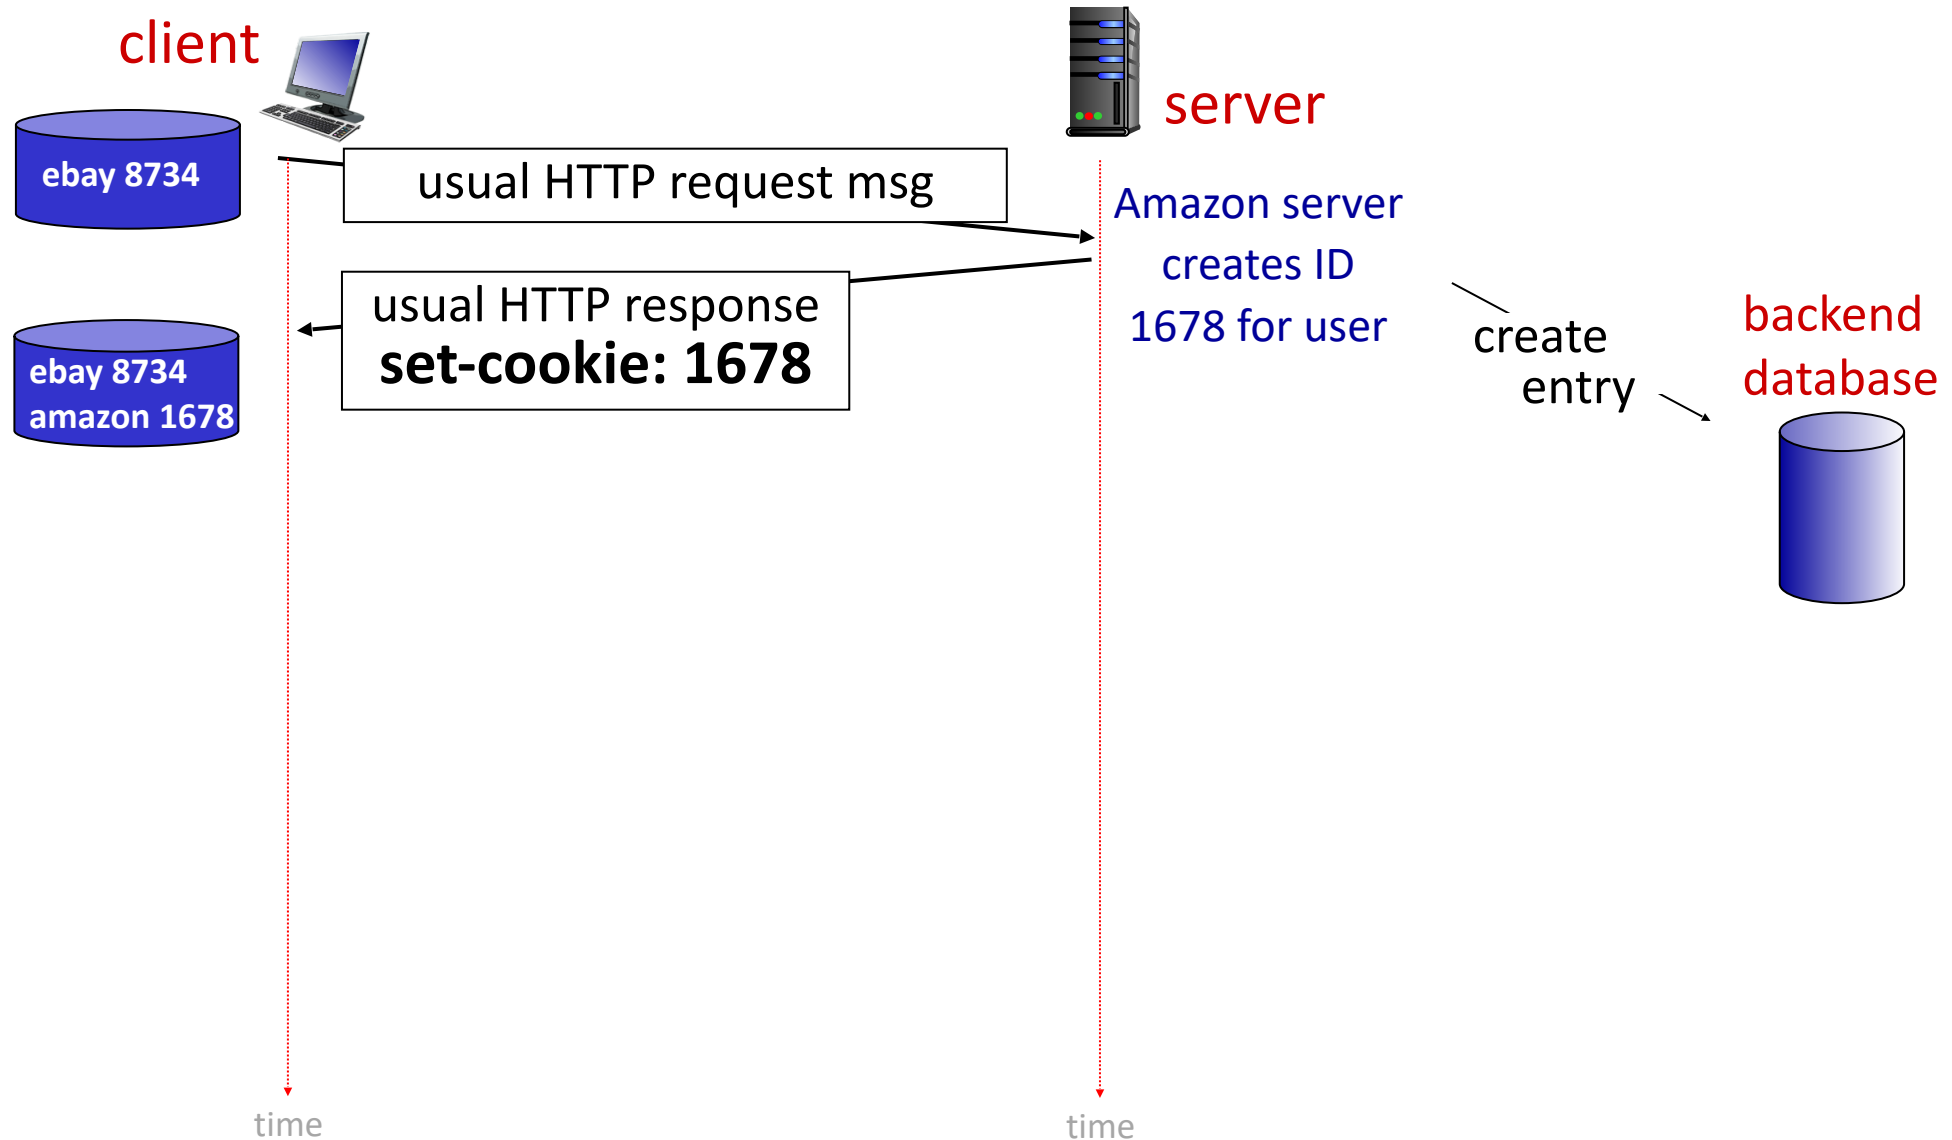
time

time

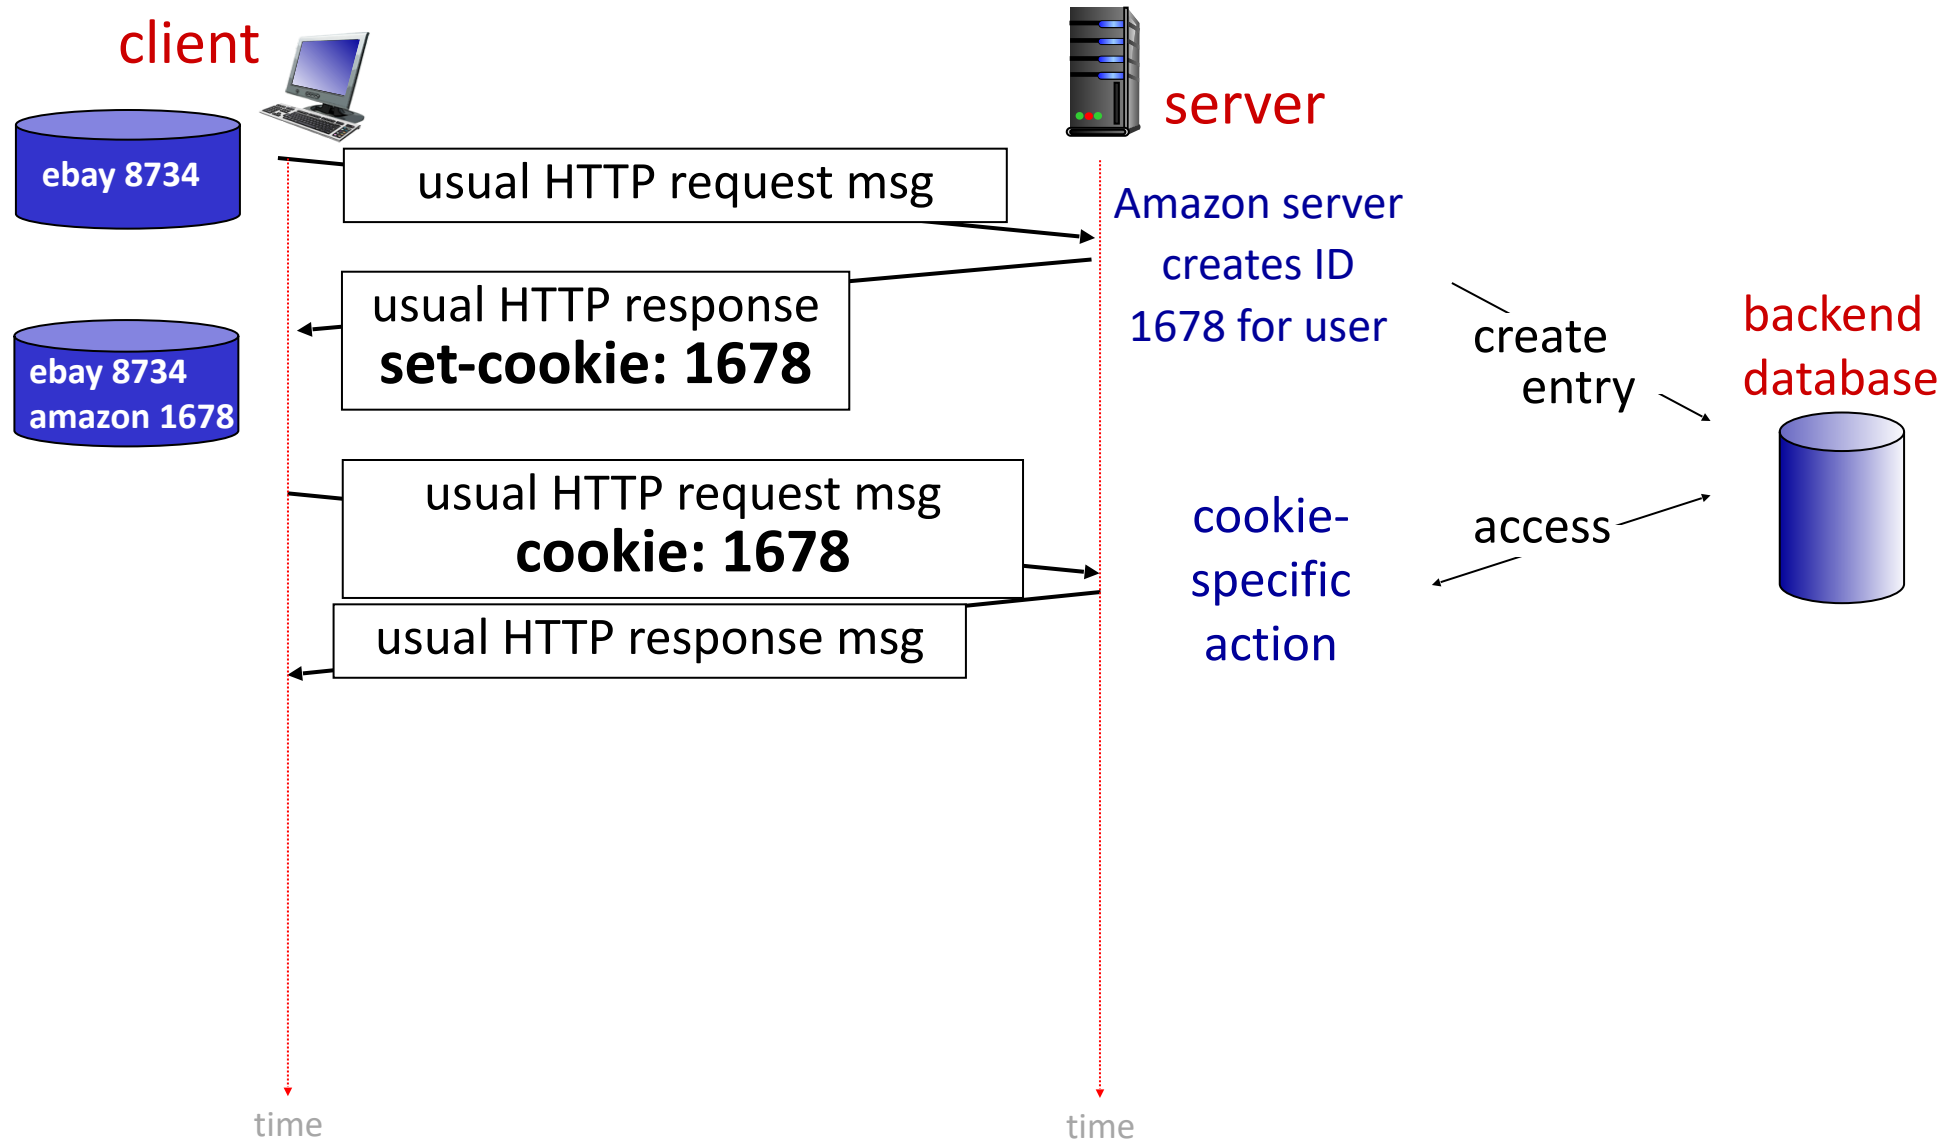
# Maintaining user/server state: cookies



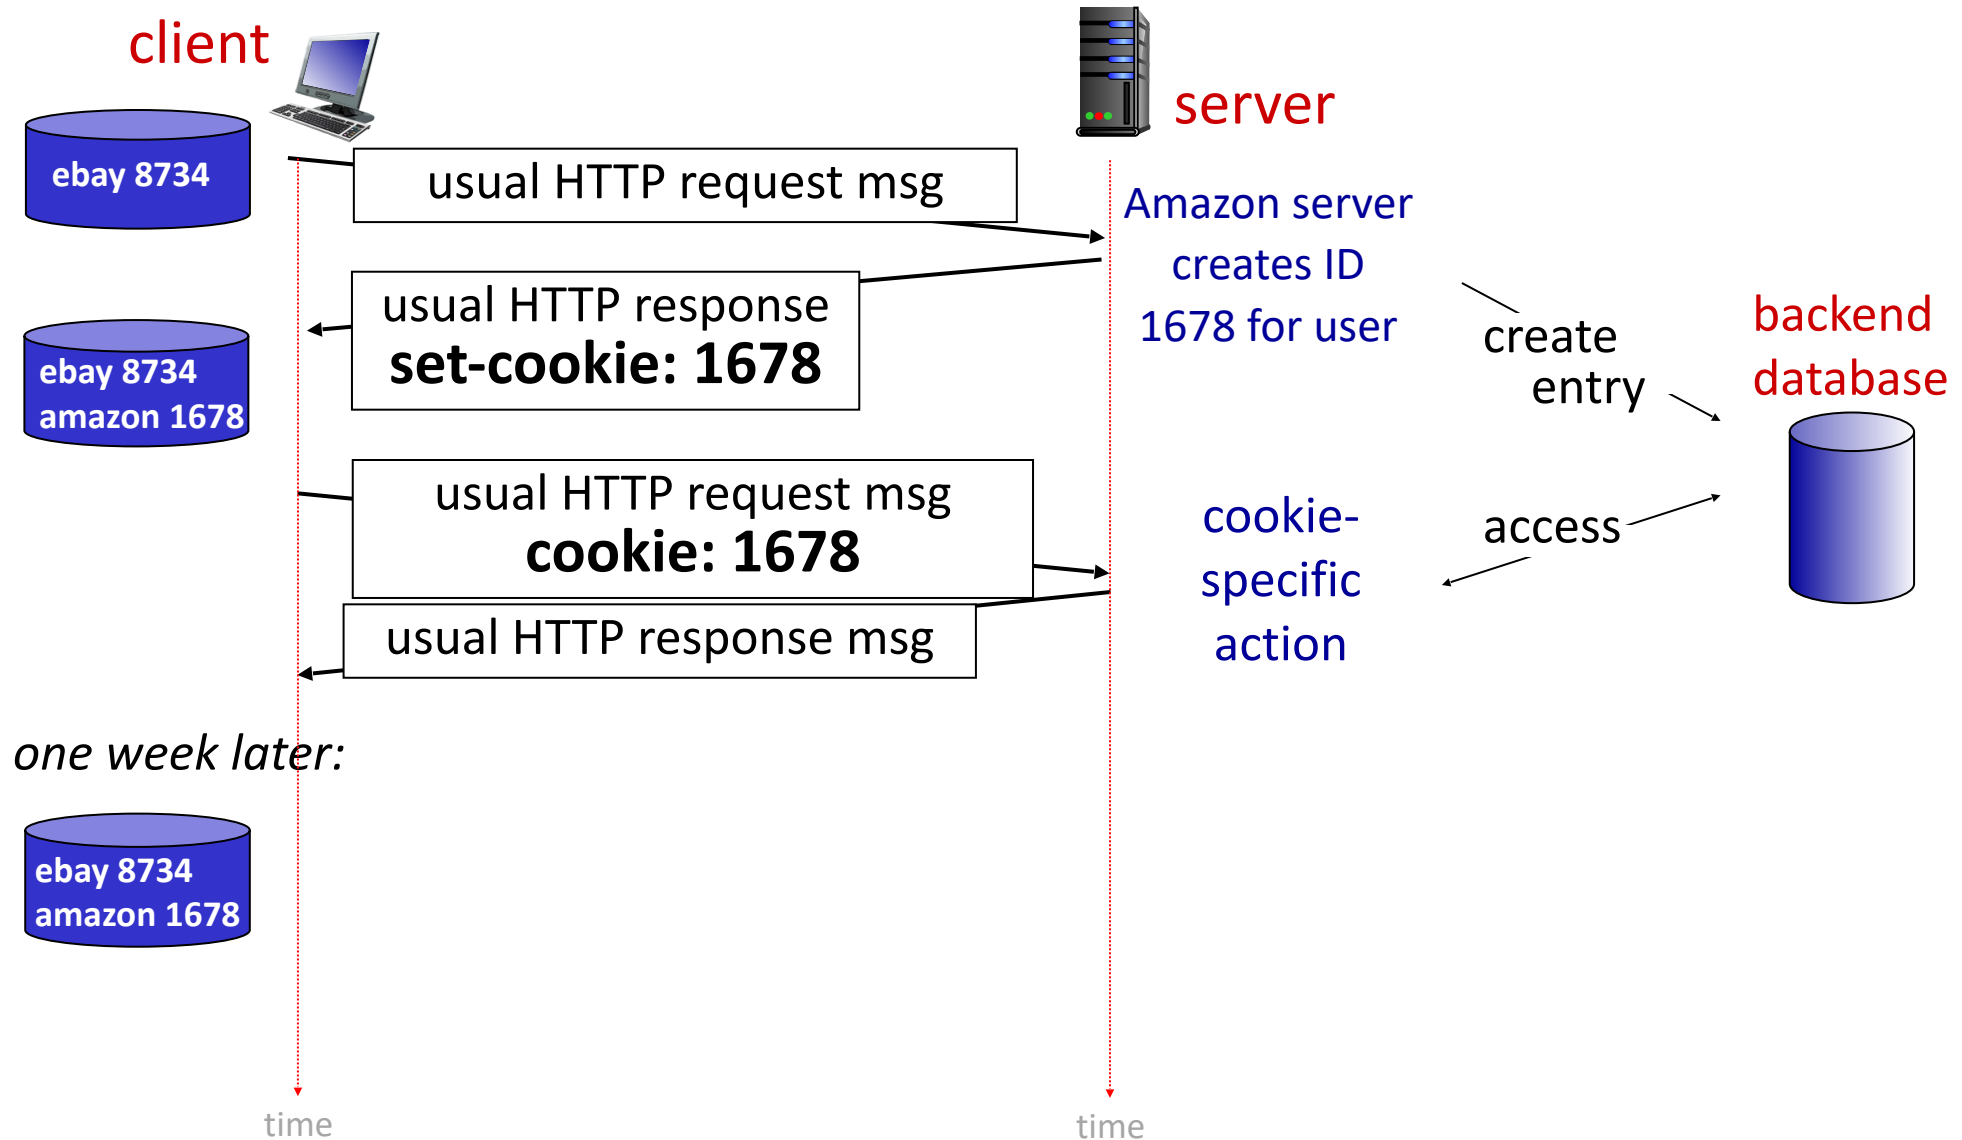
# Maintaining user/server state: cookies



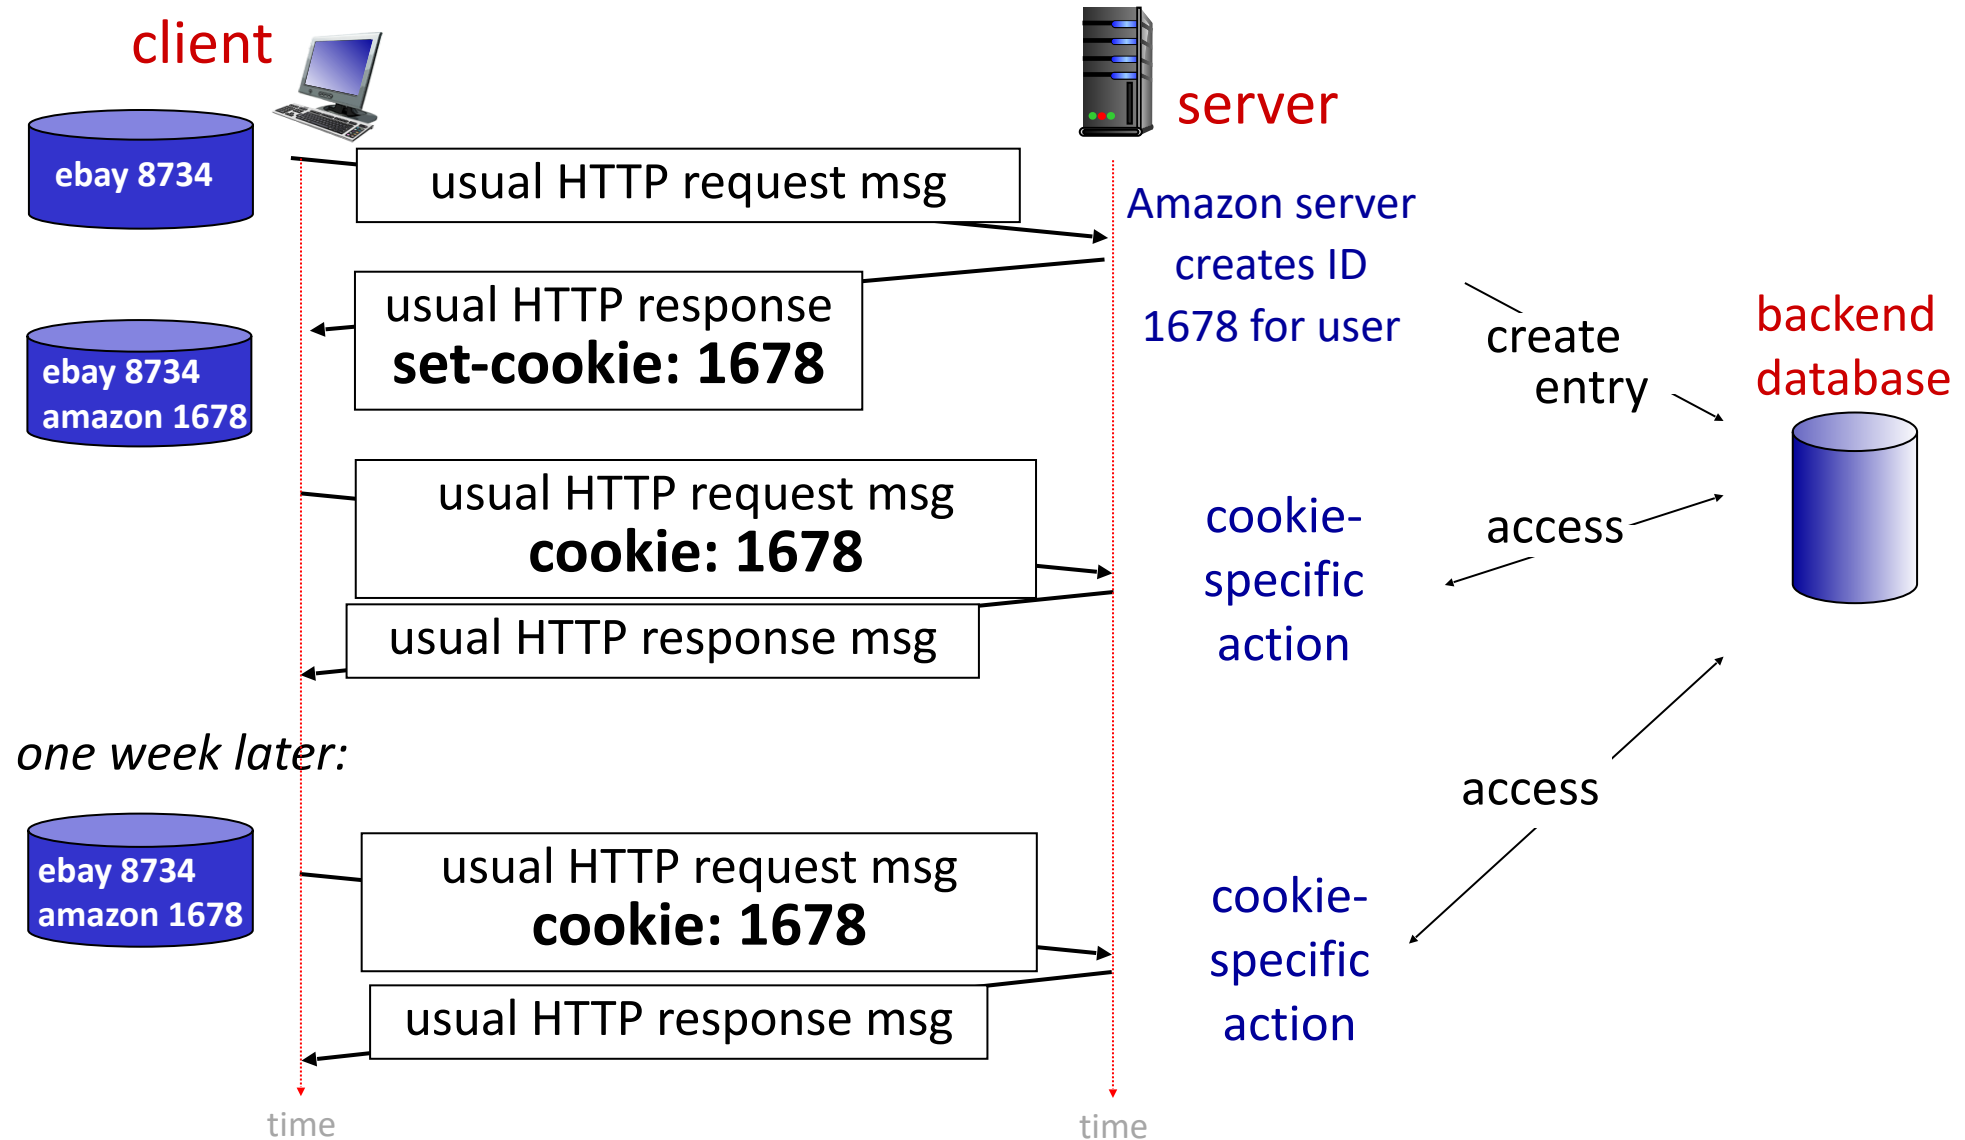
# Maintaining user/server state: cookies



# Maintaining user/server state: cookies



# Maintaining user/server state: cookies



# HTTP cookies: comments

*What cookies can be used for:*

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)



# HTTP cookies: comments

## *What cookies can be used for:*

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)

## *Challenge: How to keep state?*

- *at protocol endpoints:* maintain state at sender/receiver over multiple transactions
- *in messages:* cookies in HTTP messages carry state

# HTTP cookies: comments

## *What cookies can be used for:*

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)

## *Challenge: How to keep state?*

- *at protocol endpoints:* maintain state at sender/receiver over multiple transactions
- *in messages:* cookies in HTTP messages carry state

— aside —  
*cookies and privacy:*

- cookies permit sites to *learn* a lot about you on their site.
- third party persistent cookies (tracking cookies) allow common identity (cookie value) to be tracked across multiple web sites

# Resumo de HTTP

HTTP (Cookies)

Proxy

- ☐ Modelo Cliente/Servidor
- ☐ Requisições do cliente em ASCII
- ☐ Protocolo sem estado se não utilizar cookies
- ☐ Persistente X Não-persistente
- ☐ Em termos de sockets
  - O cliente envia o primeiro `write`
  - O servidor pode manter a mesma conexão TCP aberta

HTTP (Cookies)

▶ Proxy

Proxy

# Web caches

*Goal:* satisfy client requests without involving origin server



client



origin  
server



client

# Web caches

*Goal:* satisfy client requests without involving origin server

- user configures browser to point to a (local) *Web cache*



client

Web  
cache



origin  
server



client

# Web caches

*Goal:* satisfy client requests without involving origin server

- user configures browser to point to a (local) *Web cache*
- browser sends all HTTP requests to cache
  - *if* object in cache: cache returns object to client
  - *else* cache requests object from origin server, caches received object, then returns object to client



client

Web  
cache



origin  
server

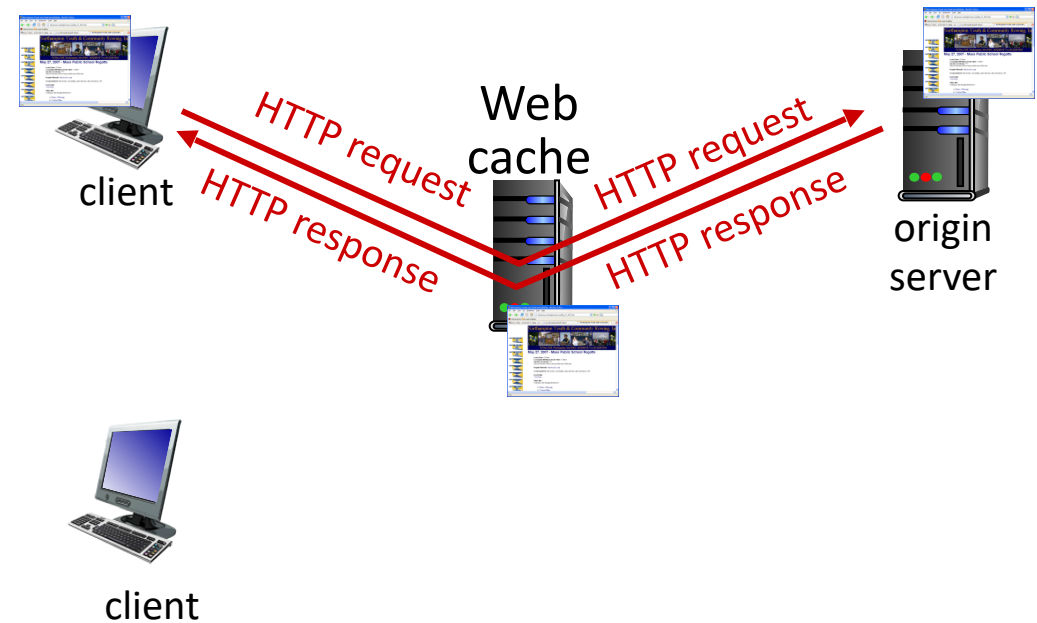


client

# Web caches

*Goal:* satisfy client requests without involving origin server

- user configures browser to point to a (local) *Web cache*
- browser sends all HTTP requests to cache
  - *if* object in cache: cache returns object to client
  - *else* cache requests object from origin server, caches received object, then returns object to client

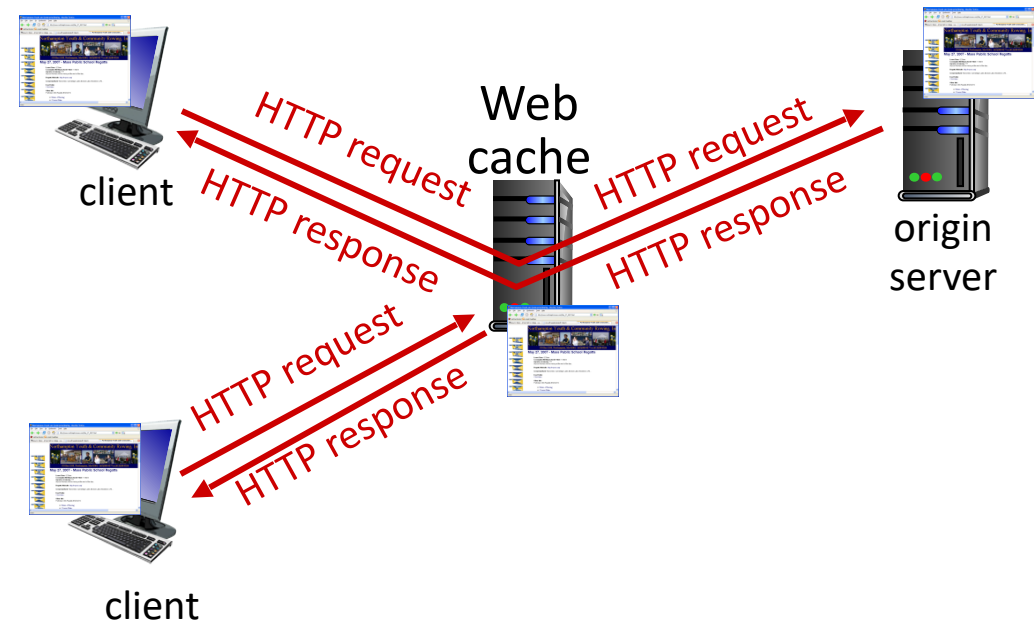




# Web caches

*Goal:* satisfy client requests without involving origin server

- user configures browser to point to a (local) *Web cache*
- browser sends all HTTP requests to cache
  - *if* object in cache: cache returns object to client
  - *else* cache requests object from origin server, caches received object, then returns object to client



# Web caches (aka proxy servers)

- Web cache acts as both client and server
  - server for original requesting client
  - client to origin server

# Web caches (aka proxy servers)

- Web cache acts as both client and server
  - server for original requesting client
  - client to origin server
- server tells cache about object's allowable caching in response header:

```
Cache-Control: max-age=<seconds>
```

```
Cache-Control: no-cache
```

# Web caches (aka proxy servers)

- Web cache acts as both client and server
  - server for original requesting client
  - client to origin server
- server tells cache about object's allowable caching in response header:

```
Cache-Control: max-age=<seconds>
```

```
Cache-Control: no-cache
```

*Why* Web caching?

# Web caches (aka proxy servers)

- Web cache acts as both client and server
  - server for original requesting client
  - client to origin server
- server tells cache about object's allowable caching in response header:

```
Cache-Control: max-age=<seconds>
```

```
Cache-Control: no-cache
```

## *Why* Web caching?

- reduce response time for client request
  - cache is closer to client

# Web caches (aka proxy servers)

- Web cache acts as both client and server
  - server for original requesting client
  - client to origin server
- server tells cache about object's allowable caching in response header:

```
Cache-Control: max-age=<seconds>
```

```
Cache-Control: no-cache
```

## *Why* Web caching?

- reduce response time for client request
  - cache is closer to client
- reduce traffic on an institution's access link

# Web caches (aka proxy servers)

- Web cache acts as both client and server
  - server for original requesting client
  - client to origin server
- server tells cache about object's allowable caching in response header:

```
Cache-Control: max-age=<seconds>
```

```
Cache-Control: no-cache
```

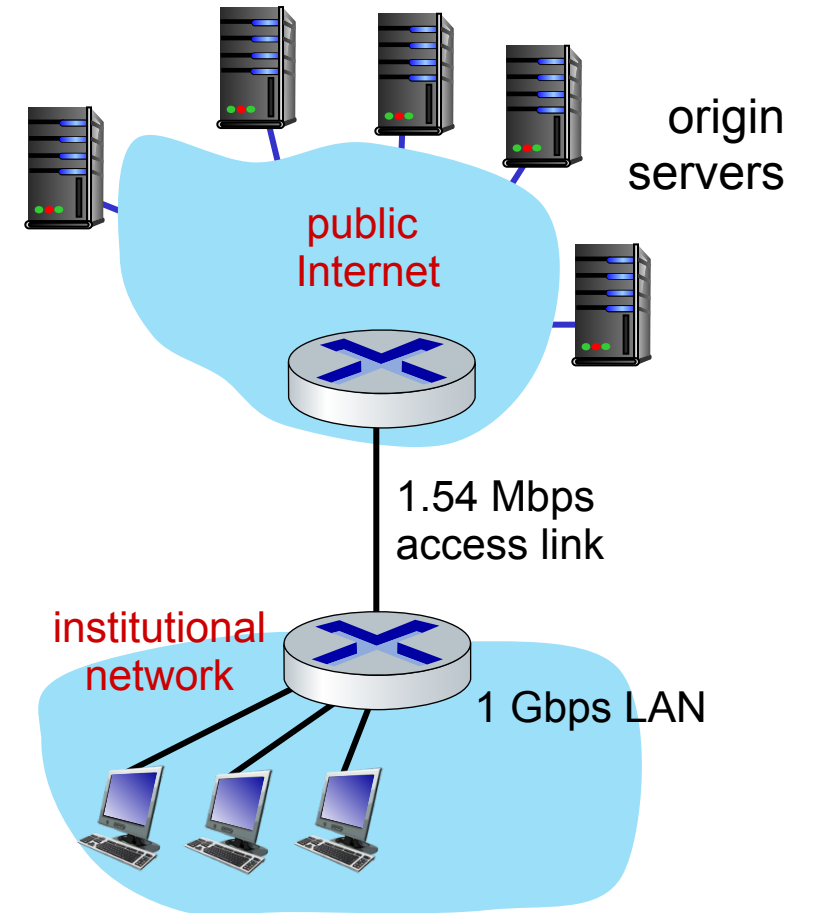
## *Why* Web caching?

- reduce response time for client request
  - cache is closer to client
- reduce traffic on an institution's access link
- Internet is dense with caches
  - enables “poor” content providers to more effectively deliver content

# Caching example

## *Scenario:*

- access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
  - avg data rate to browsers: 1.50 Mbps





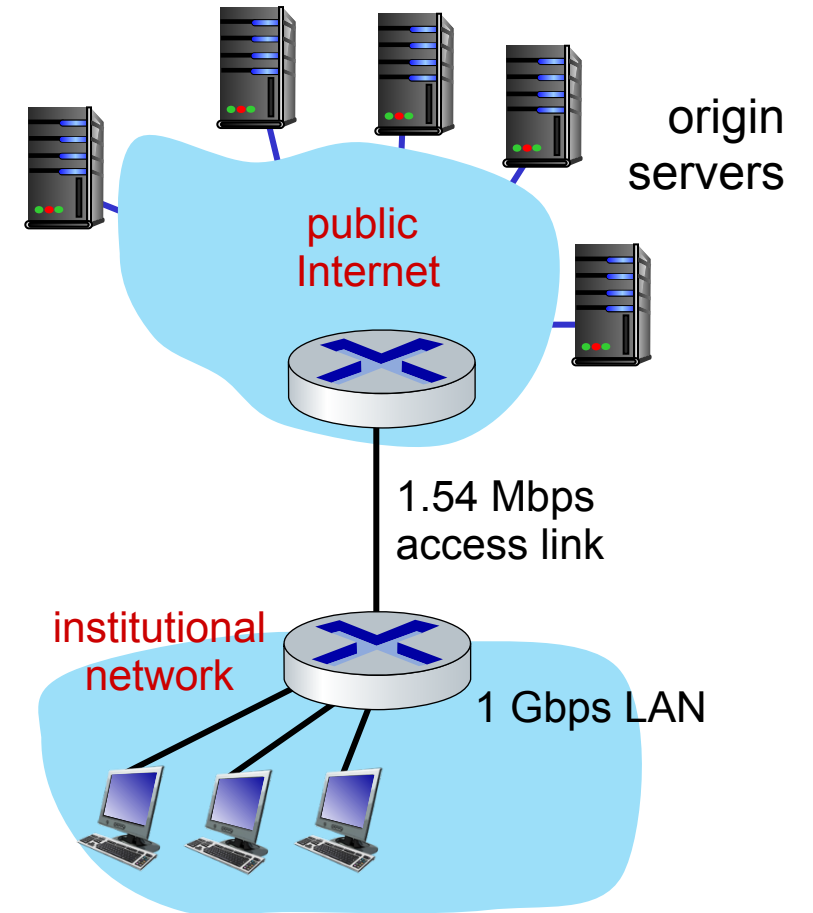
# Caching example

## *Scenario:*

- access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
  - avg data rate to browsers: 1.50 Mbps

## *Performance:*

- access link utilization = .97



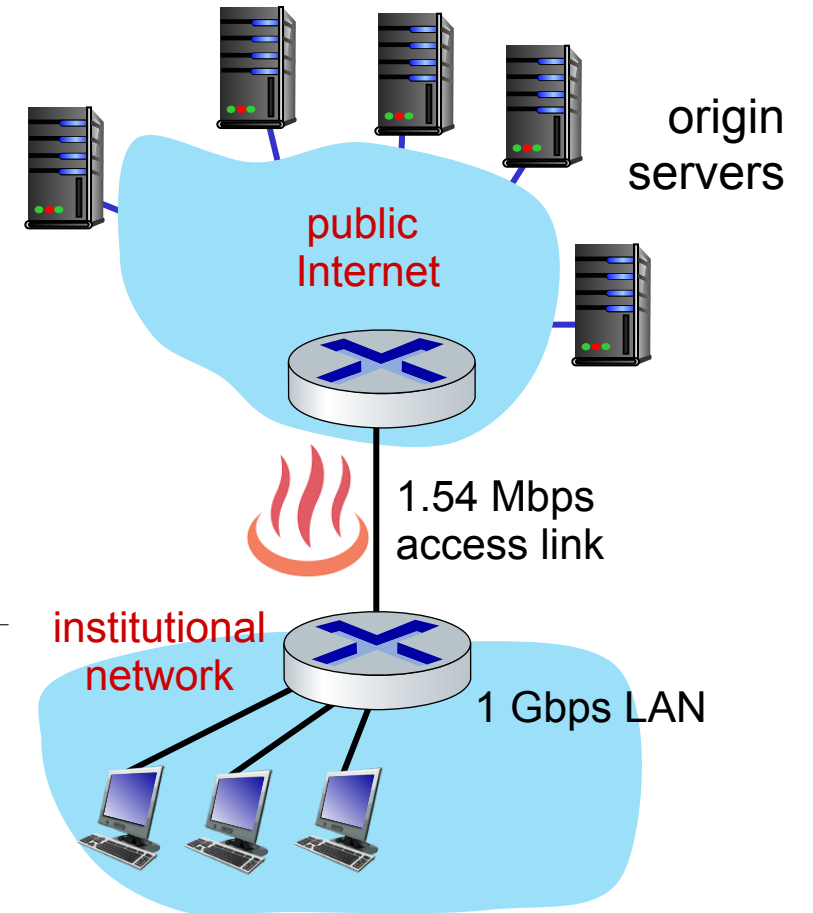
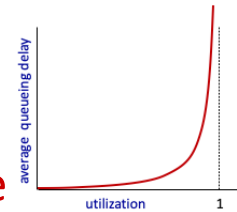
# Caching example

## Scenario:

- access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
  - avg data rate to browsers: 1.50 Mbps

## Performance:

- access link utilization = .97
- problem: large queueing delays at high utilization!*



# Caching example

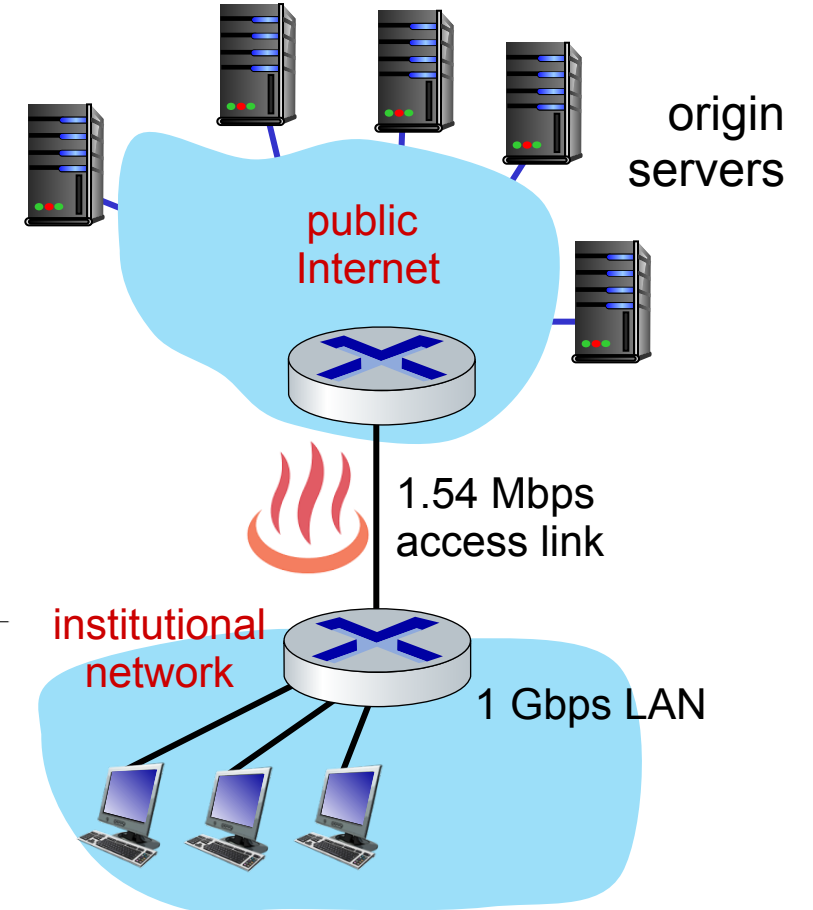
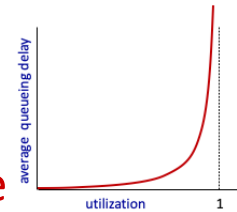
## Scenario:

- access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
  - avg data rate to browsers: 1.50 Mbps

## Performance:

- access link utilization = .97
- LAN utilization: .0015

*problem: large queueing delays at high utilization!*



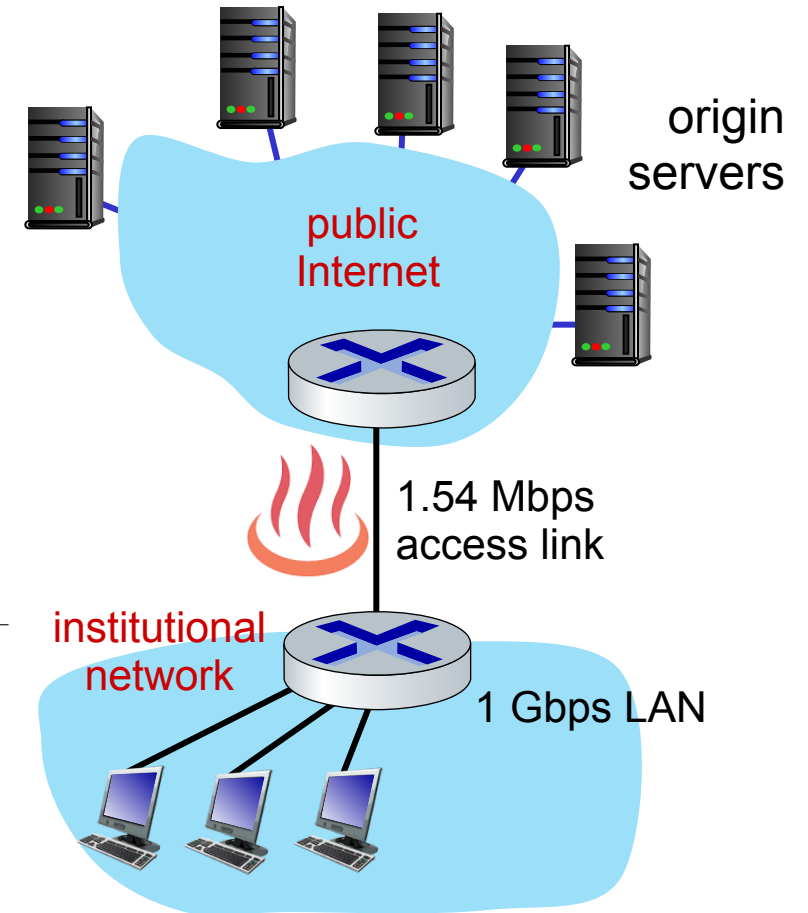
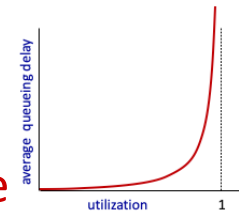
# Caching example

## Scenario:

- access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
  - avg data rate to browsers: 1.50 Mbps

## Performance:

- access link utilization = **.97** *problem: large queueing delays at high utilization!*
- LAN utilization: .0015
- end-end delay = Internet delay +  
access link delay + LAN delay  
= 2 sec + minutes + usecs



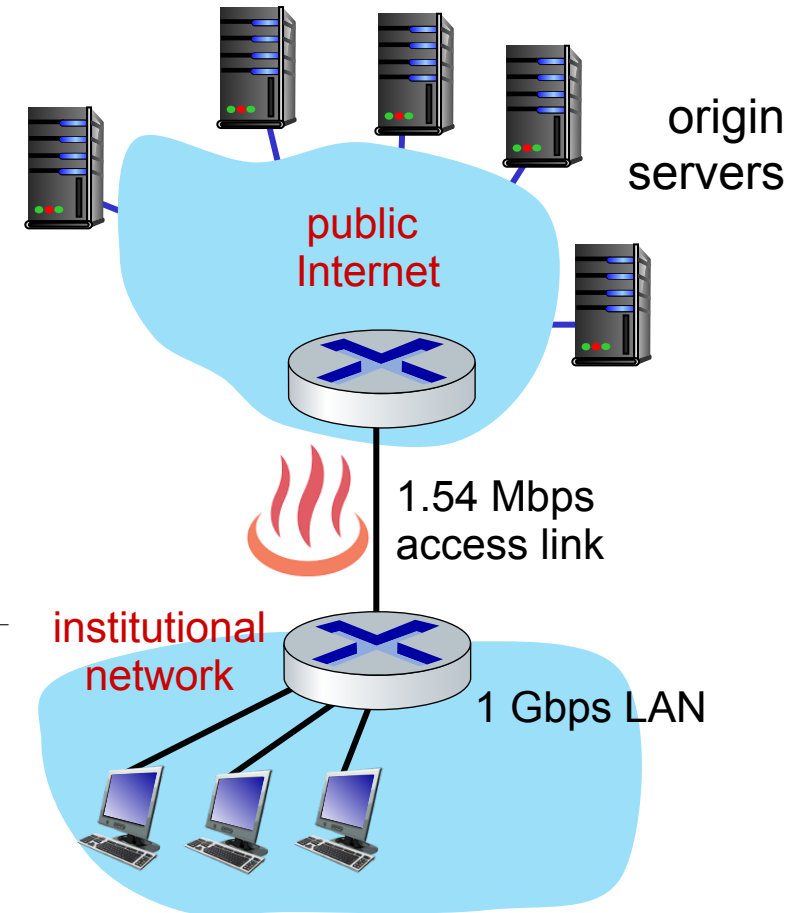
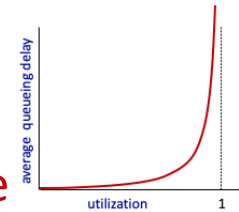
# Caching example

## Scenario:

- access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
  - avg data rate to browsers: 1.50 Mbps

## Performance:

- access link utilization = **.97** *problem: large queueing delays at high utilization!*
- LAN utilization: .0015
- end-end delay = Internet delay + access link delay + LAN delay  
= 2 sec + **minutes** + usecs



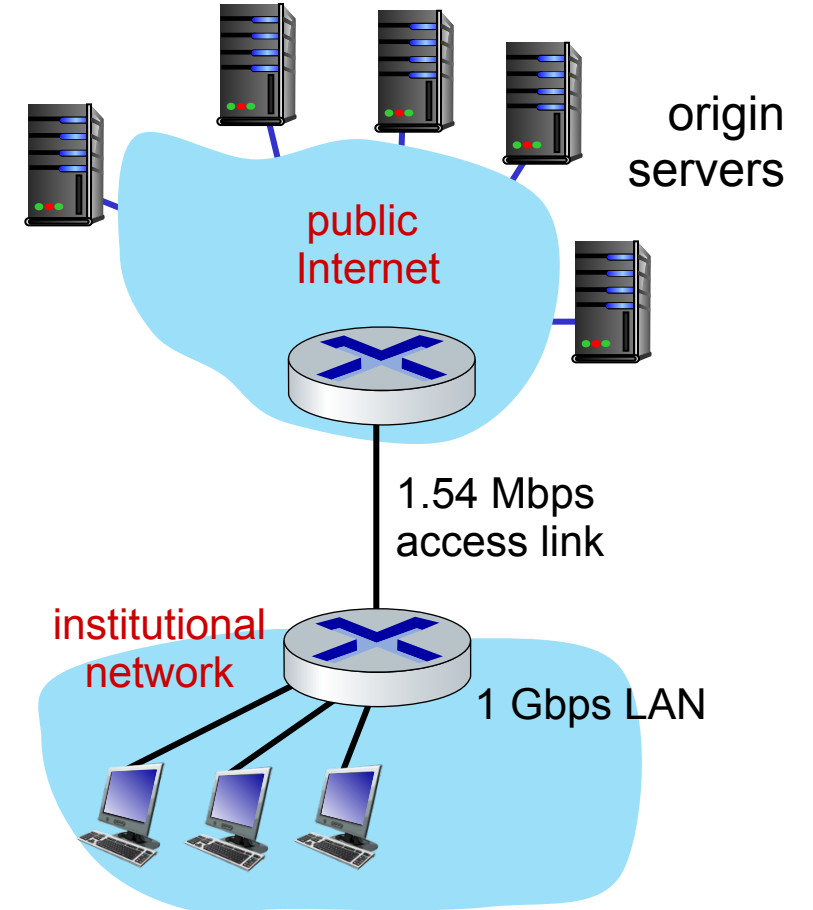
# Option 1: buy a faster access link

## Scenario:

- access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
  - avg data rate to browsers: 1.50 Mbps

## Performance:

- access link utilization = .97
- LAN utilization: .0015
- end-end delay = Internet delay +  
access link delay + LAN delay  
= 2 sec + minutes + usecs



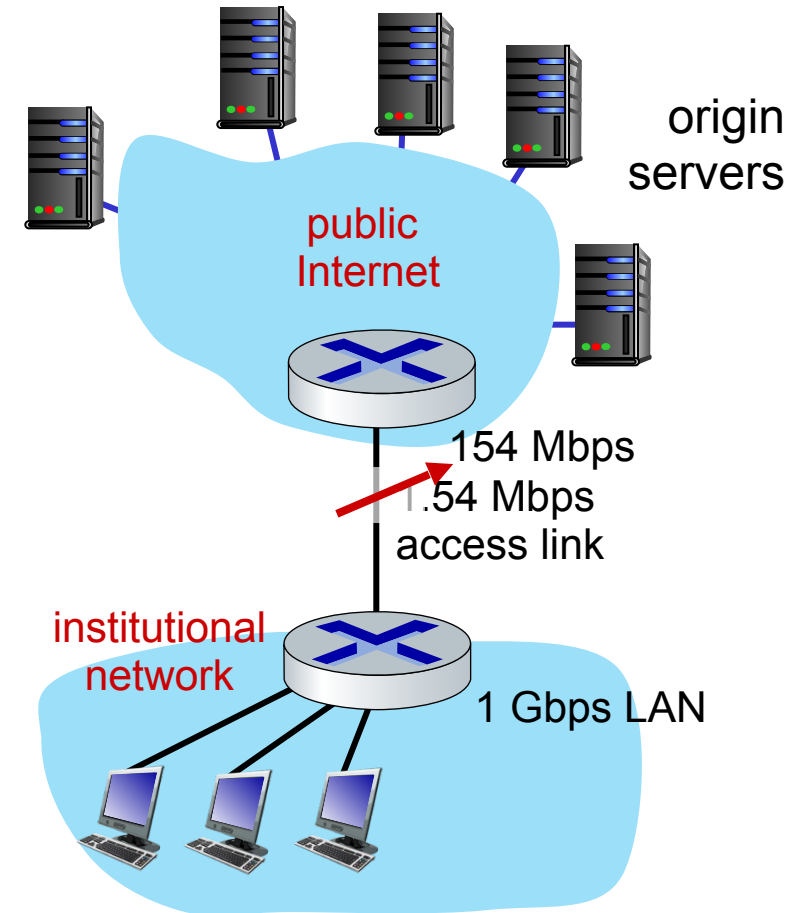
# Option 1: buy a faster access link

## Scenario:

- access link rate: ~~1.54~~ 154 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
  - avg data rate to browsers: 1.50 Mbps

## Performance:

- access link utilization = .97
- LAN utilization: .0015
- end-end delay = Internet delay +  
access link delay + LAN delay  
= 2 sec + minutes + usecs



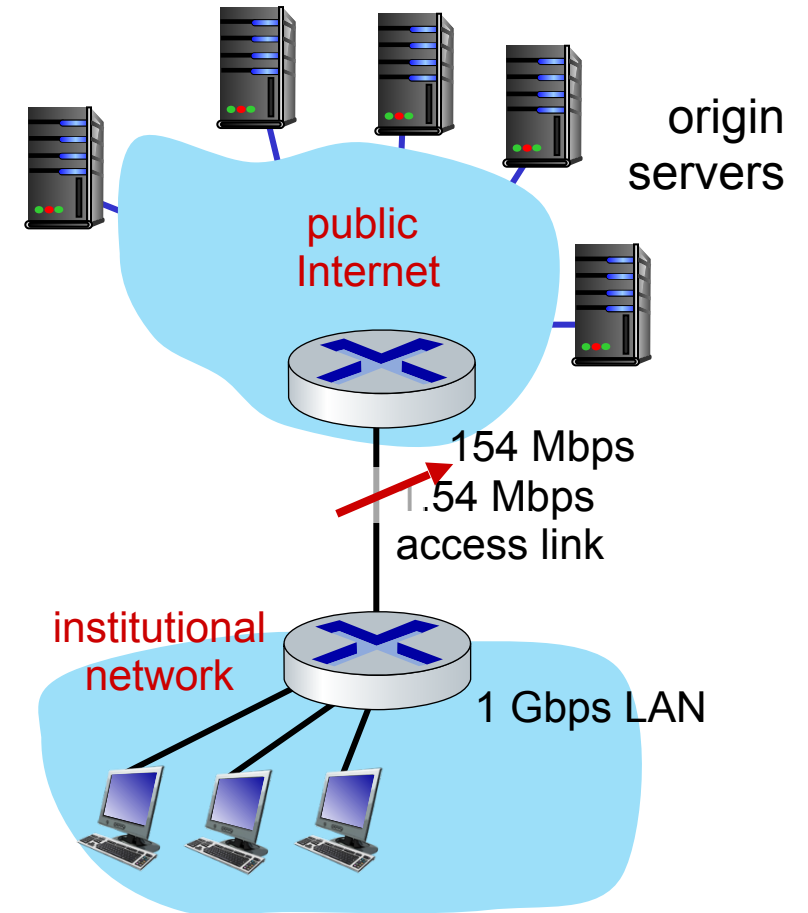
# Option 1: buy a faster access link

## Scenario:

- access link rate: ~~1.54~~ 154 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
  - avg data rate to browsers: 1.50 Mbps

## Performance:

- access link utilization = ~~.97~~ → .0097
- LAN utilization: .0015
- end-end delay = Internet delay +  
access link delay + LAN delay  
= 2 sec + minutes + usecs





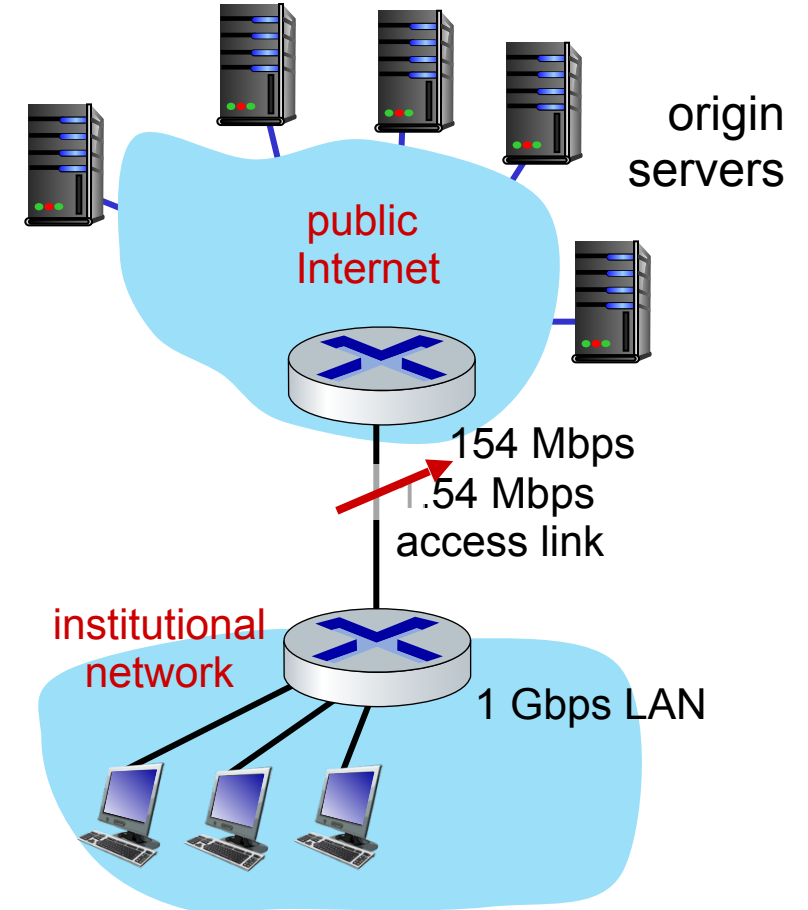
# Option 1: buy a faster access link

## Scenario:

- access link rate: ~~1.54~~ 154 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
  - avg data rate to browsers: 1.50 Mbps

## Performance:

- access link utilization = ~~.97~~ → .0097
- LAN utilization: .0015
- end-end delay = Internet delay +  
access link delay + LAN delay  
= 2 sec + ~~minutes~~ → msec



# Option 1: buy a faster access link

## Scenario:

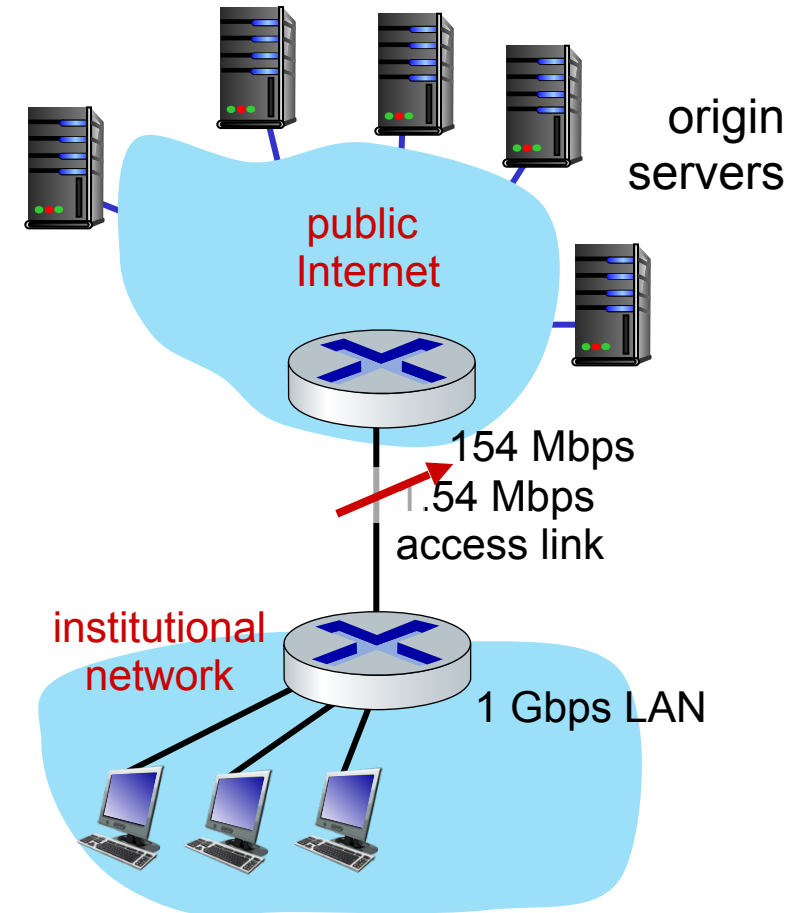
- access link rate: ~~1.54~~ 154 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
  - avg data rate to browsers: 1.50 Mbps

## Performance:

- access link utilization = ~~.97~~ .0097
- LAN utilization: .0015
- end-end delay = Internet delay +  
access link delay + LAN delay

$$= 2 \text{ sec} + \text{minutes} + \text{usecs}$$

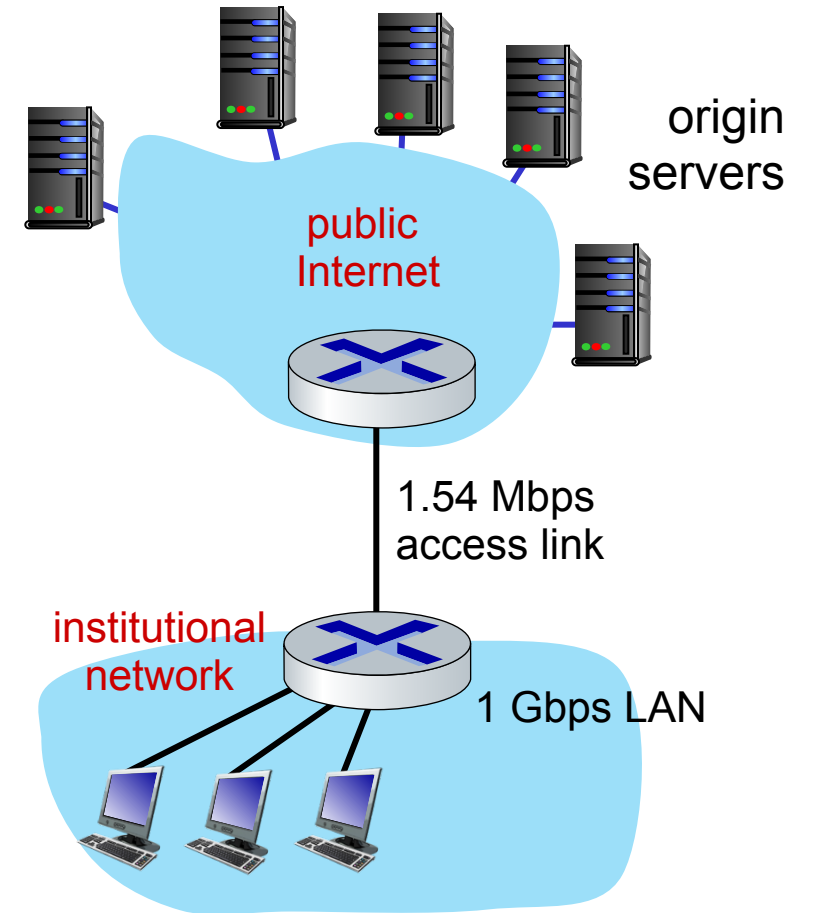
**Cost:** faster access link (expensive!) ~~msecs~~



# Option 2: install a web cache

## *Scenario:*

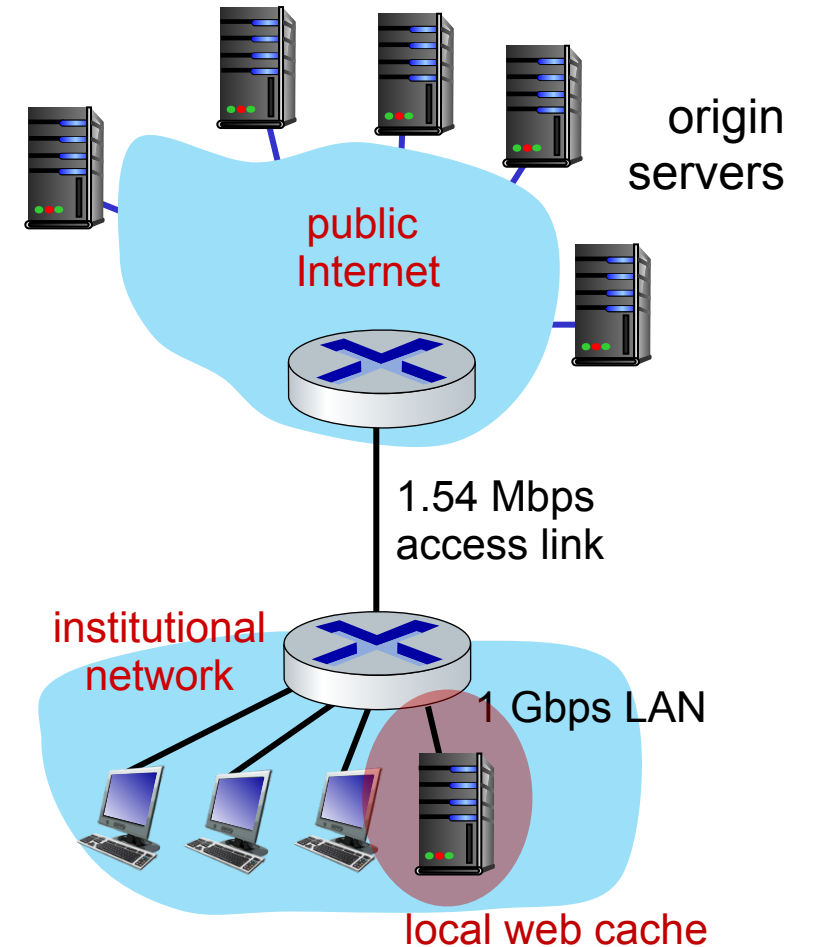
- access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
  - avg data rate to browsers: 1.50 Mbps



# Option 2: install a web cache

## *Scenario:*

- access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
  - avg data rate to browsers: 1.50 Mbps

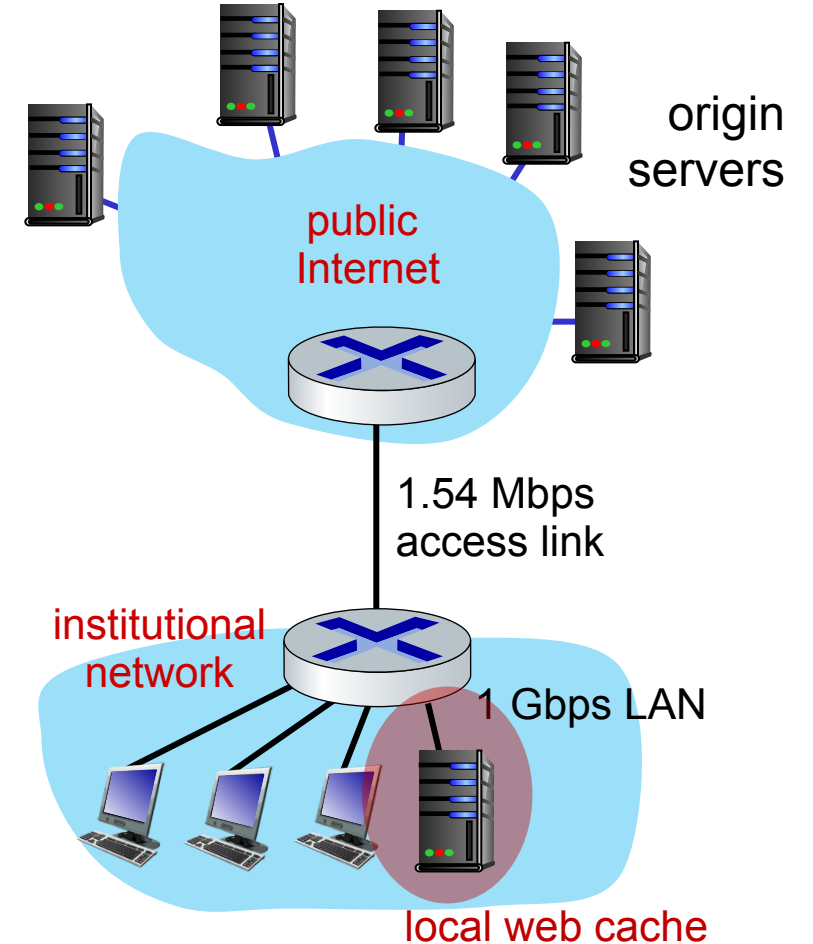


# Option 2: install a web cache

## *Scenario:*

- access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
  - avg data rate to browsers: 1.50 Mbps

*Cost:* web cache (cheap!)



# Option 2: install a web cache

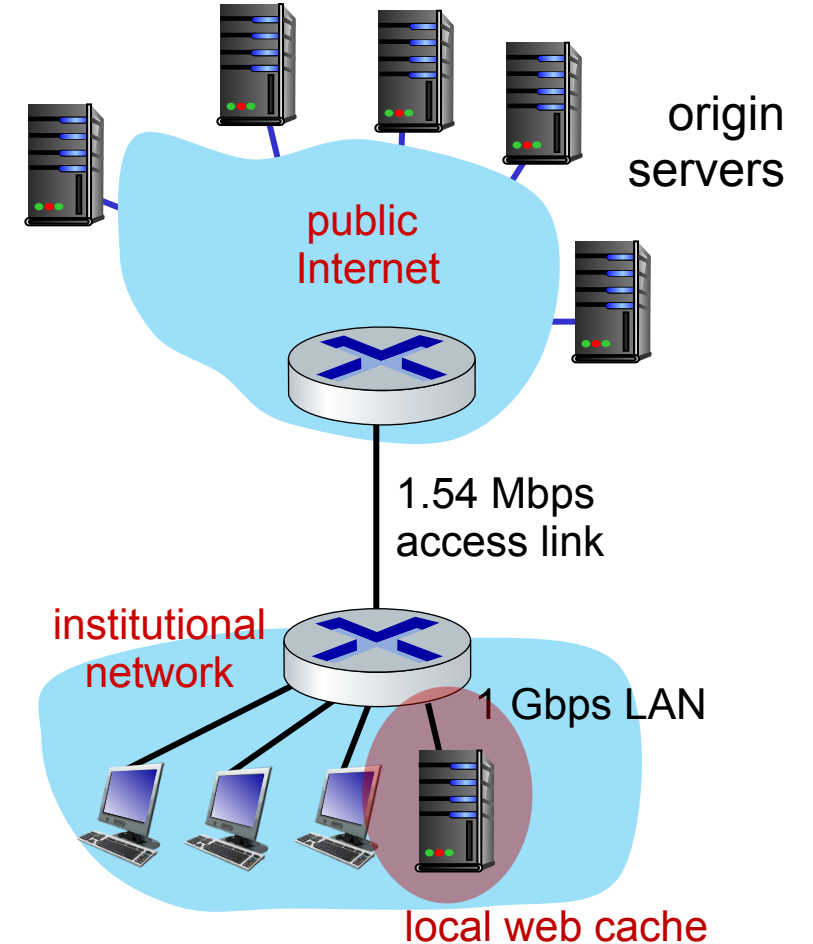
## *Scenario:*

- access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
  - avg data rate to browsers: 1.50 Mbps

*Cost:* web cache (cheap!)

## *Performance:*

- LAN utilization: .?
- access link utilization = ?
- average end-end delay = ?



# Option 2: install a web cache

## *Scenario:*

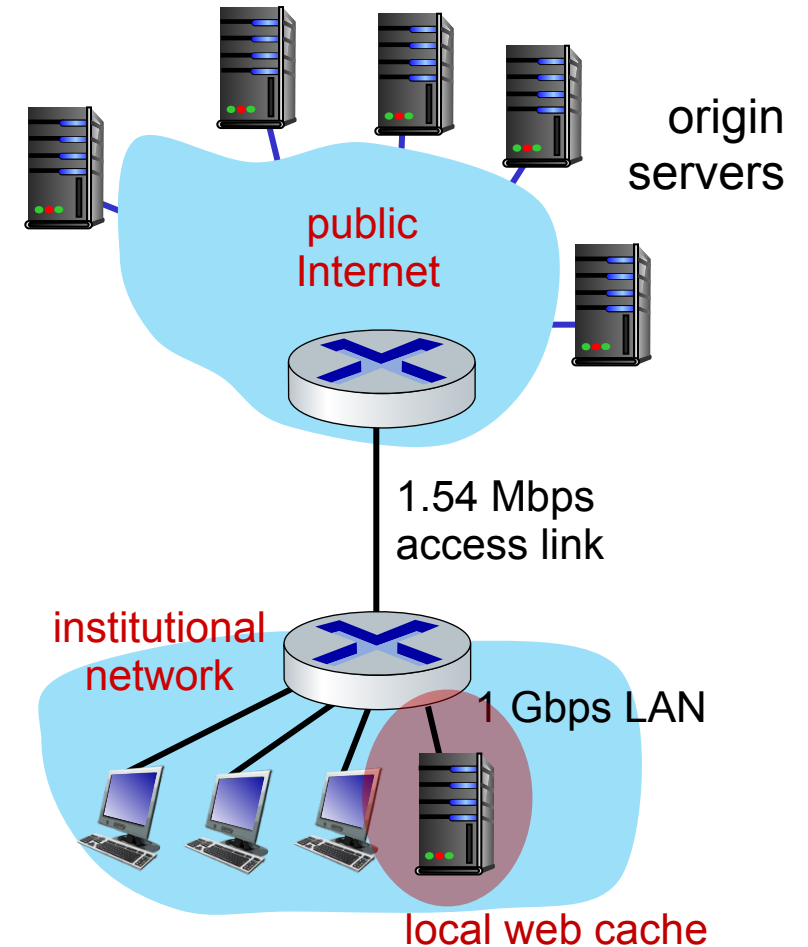
- access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
  - avg data rate to browsers: 1.50 Mbps

*Cost:* web cache (cheap!)

## *Performance:*

- LAN utilization: .?
- access link utilization = ?
- average end-end delay = ?

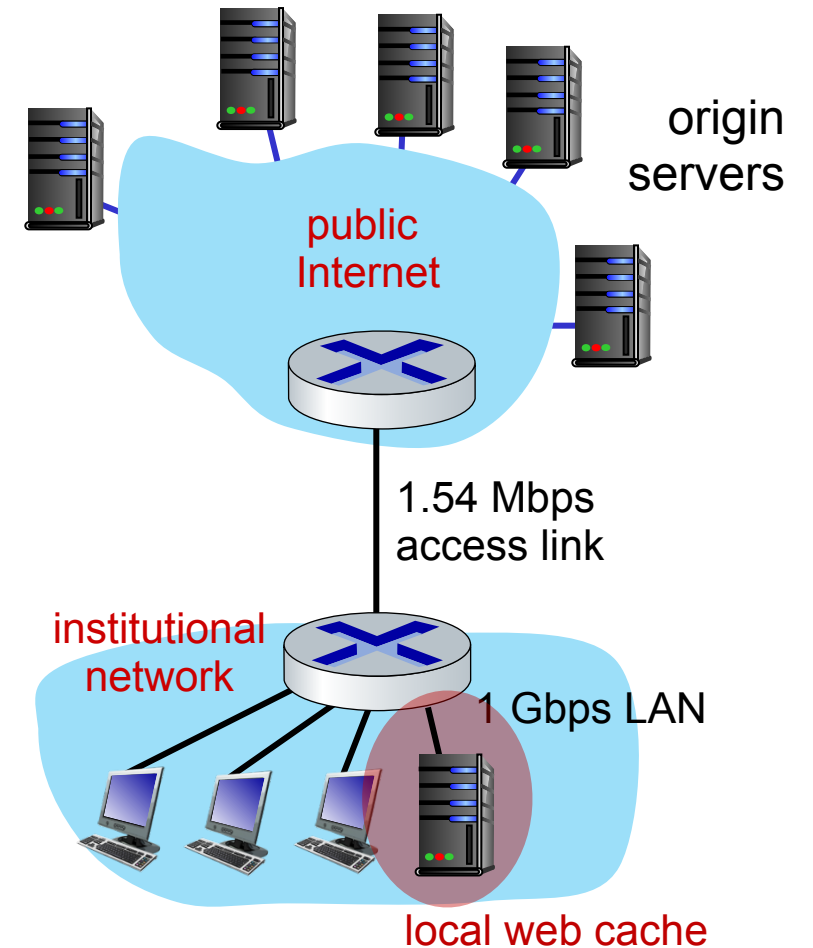
*How to compute link utilization, delay?*



# Calculating access link utilization, end-end delay with cache:

suppose cache hit rate is 0.4:

- 40% requests served by cache, with low (msec) delay

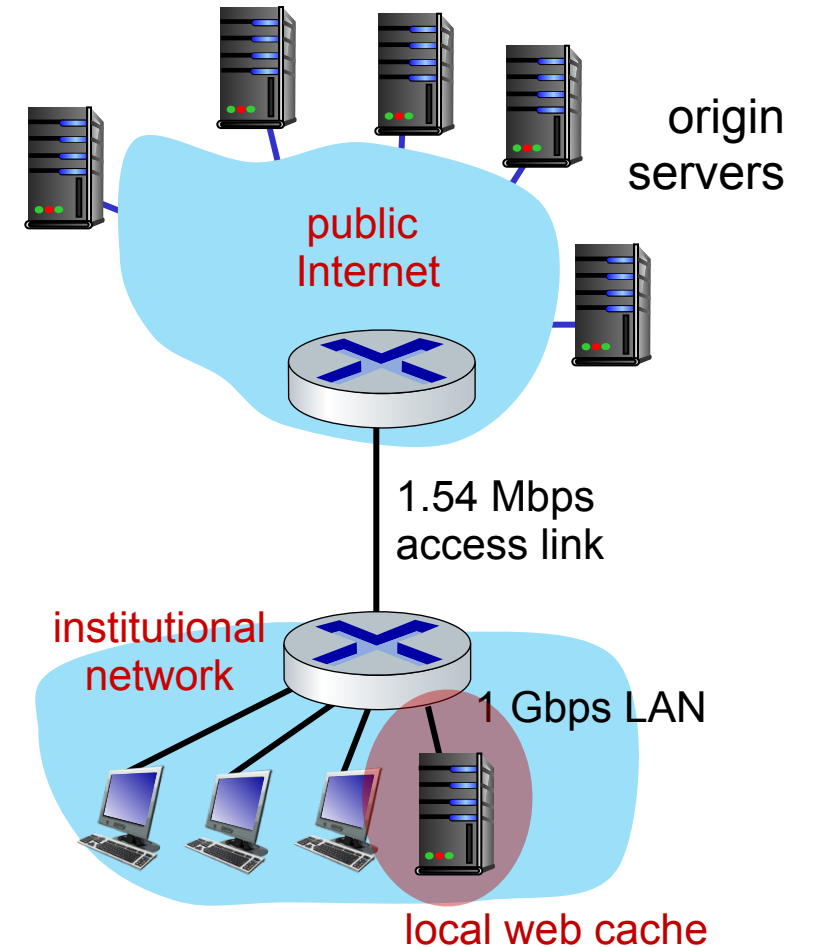




# Calculating access link utilization, end-end delay with cache:

suppose cache hit rate is 0.4:

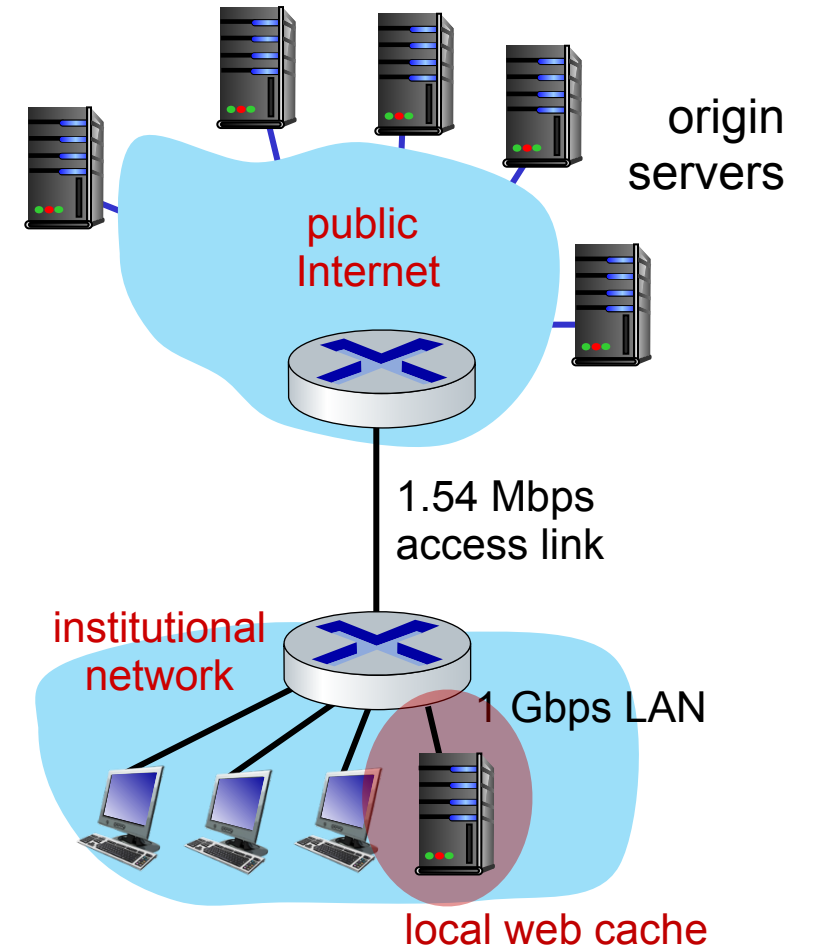
- 40% requests served by cache, with low (msec) delay
- 60% requests satisfied at origin
  - rate to browsers over access link  
 $= 0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$
  - access link utilization  $= 0.9 / 1.54 = .58$  means low (msec) queueing delay at access link



# Calculating access link utilization, end-end delay with cache:

suppose cache hit rate is 0.4:

- 40% requests served by cache, with low (msec) delay
- 60% requests satisfied at origin
  - rate to browsers over access link  
 $= 0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$
  - access link utilization  $= 0.9 / 1.54 = .58$  means low (msec) queueing delay at access link
- average end-end delay:  
 $= 0.6 * (\text{delay from origin servers})$   
 $+ 0.4 * (\text{delay when satisfied at cache})$   
 $= 0.6 (2.01) + 0.4 (\sim \text{msecs}) = \sim 1.2 \text{ secs}$

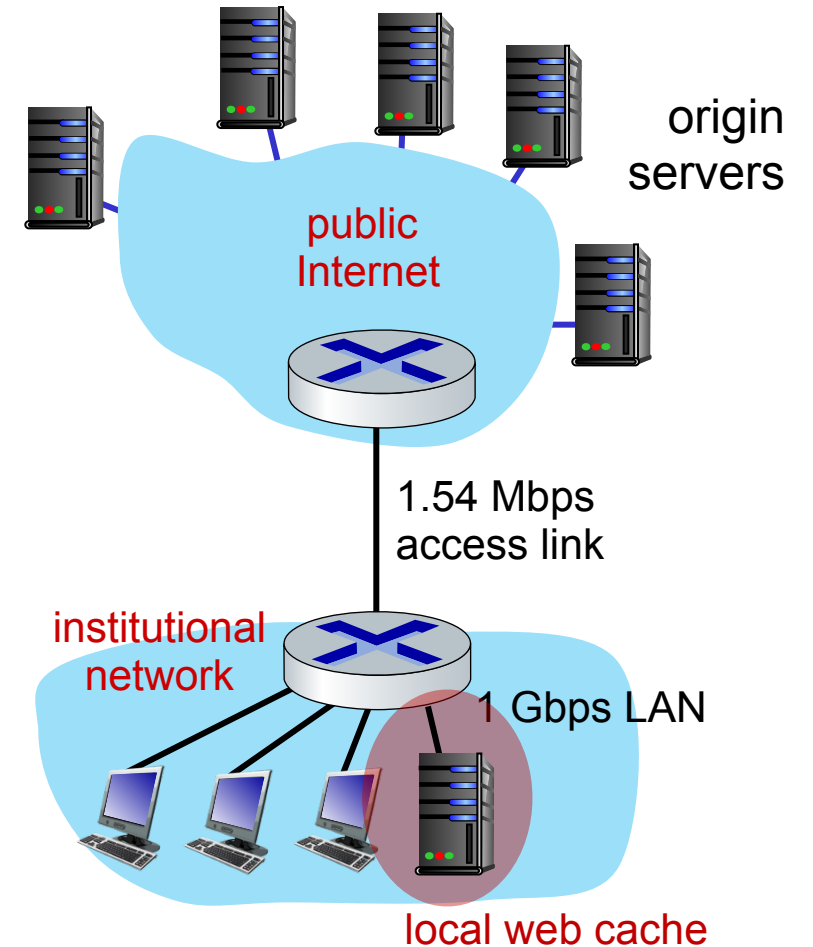


# Calculating access link utilization, end-end delay with cache:

suppose cache hit rate is 0.4:

- 40% requests served by cache, with low (msec) delay
- 60% requests satisfied at origin
  - rate to browsers over access link  
 $= 0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$
  - access link utilization  $= 0.9 / 1.54 = .58$  means low (msec) queueing delay at access link
- average end-end delay:  
 $= 0.6 * (\text{delay from origin servers})$   
 $+ 0.4 * (\text{delay when satisfied at cache})$   
 $= 0.6 (2.01) + 0.4 (\sim \text{msecs}) = \sim 1.2 \text{ secs}$

*lower average end-end delay than with 154 Mbps link (and cheaper too!)*



# Conditional GET

*Goal:* don't send object if cache has up-to-date cached version

- no object transmission delay (or use of network resources)

# Conditional GET

*Goal:* don't send object if cache has up-to-date cached version

- no object transmission delay (or use of network resources)
- *client:* specify date of cached copy in HTTP request  
**If-modified-since: <date>**

# Conditional GET

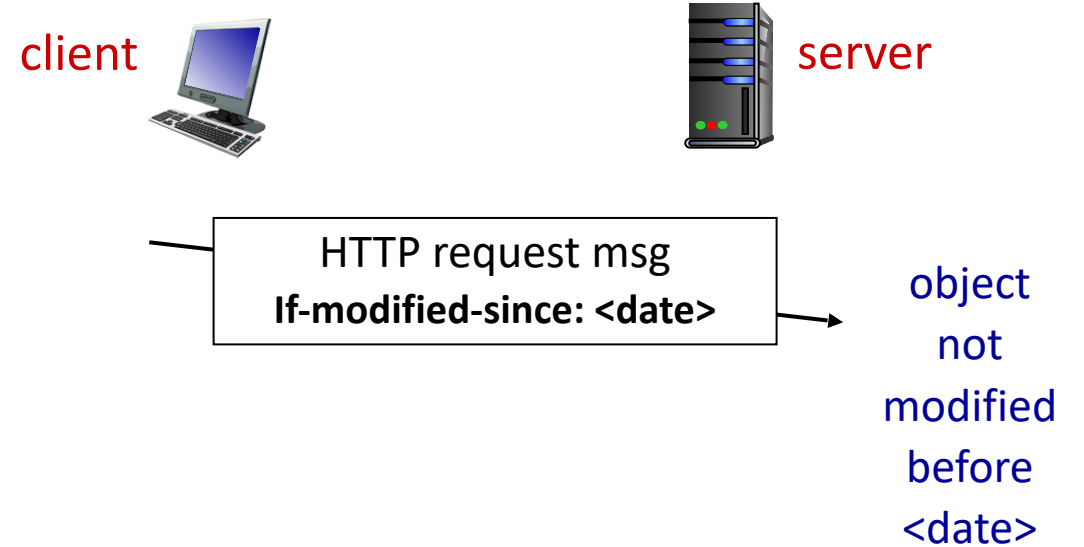
*Goal:* don't send object if cache has up-to-date cached version

- no object transmission delay (or use of network resources)
- *client:* specify date of cached copy in HTTP request  
**If-modified-since: <date>**
- *server:* response contains no object if cached copy is up-to-date:  
**HTTP/1.0 304 Not Modified**

# Conditional GET

**Goal:** don't send object if cache has up-to-date cached version

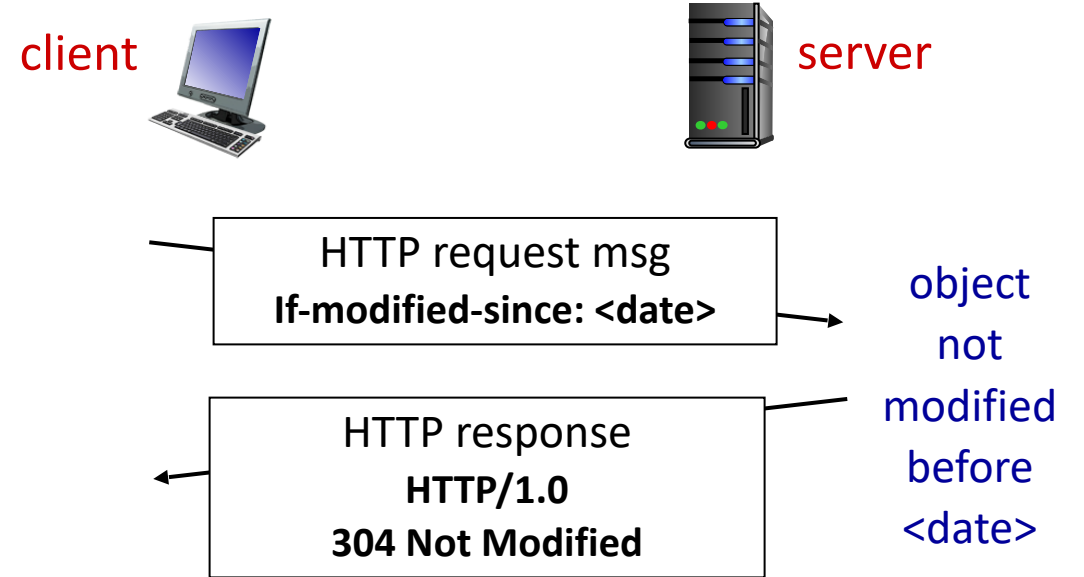
- no object transmission delay (or use of network resources)
- **client:** specify date of cached copy in HTTP request  
**If-modified-since: <date>**
- **server:** response contains no object if cached copy is up-to-date:  
**HTTP/1.0 304 Not Modified**



# Conditional GET

**Goal:** don't send object if cache has up-to-date cached version

- no object transmission delay (or use of network resources)
- **client:** specify date of cached copy in HTTP request  
**If-modified-since: <date>**
- **server:** response contains no object if cached copy is up-to-date:  
**HTTP/1.0 304 Not Modified**





# Conditional GET

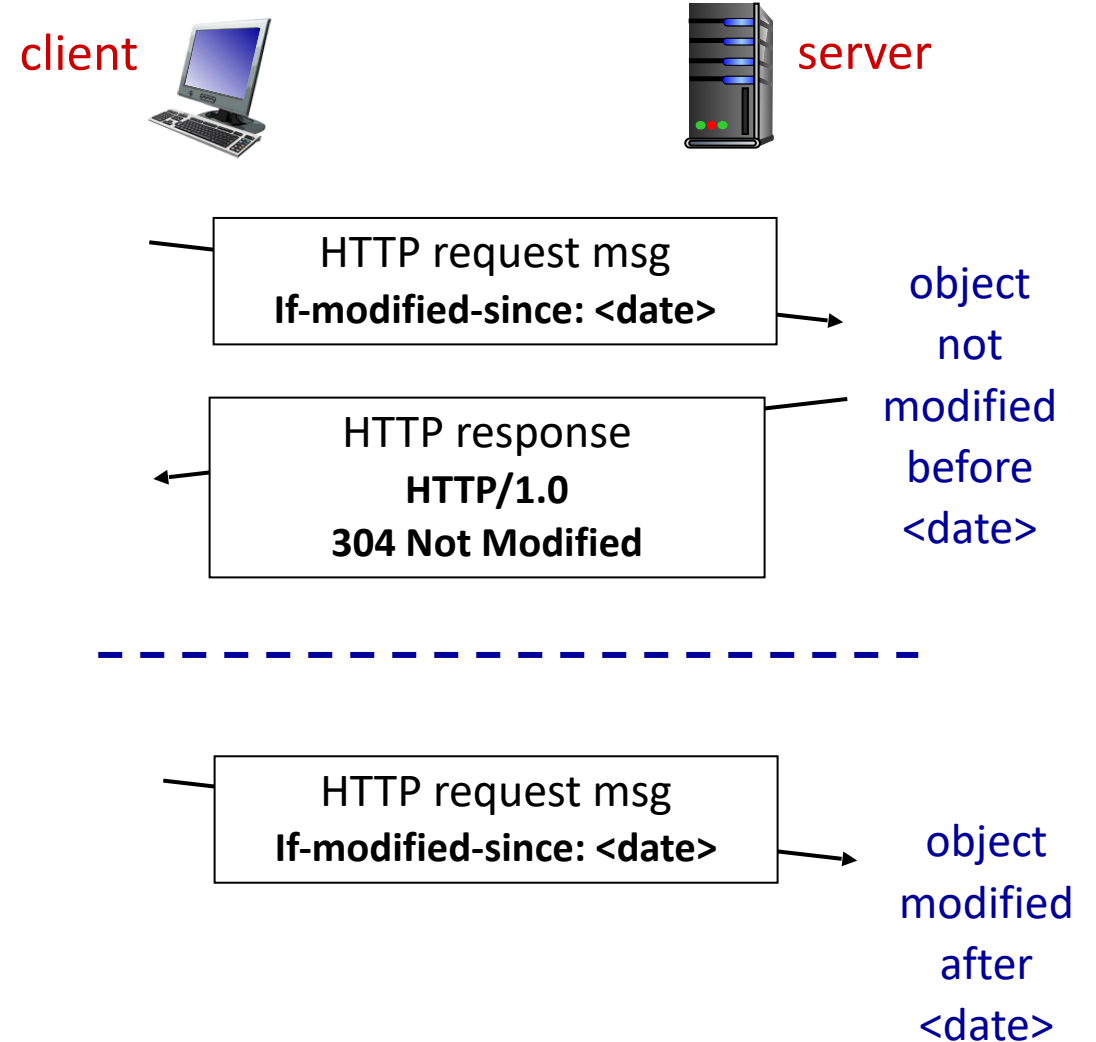
**Goal:** don't send object if cache has up-to-date cached version

- no object transmission delay (or use of network resources)

- **client:** specify date of cached copy in HTTP request

**If-modified-since: <date>**

- **server:** response contains no object if cached copy is up-to-date:  
**HTTP/1.0 304 Not Modified**



# Conditional GET

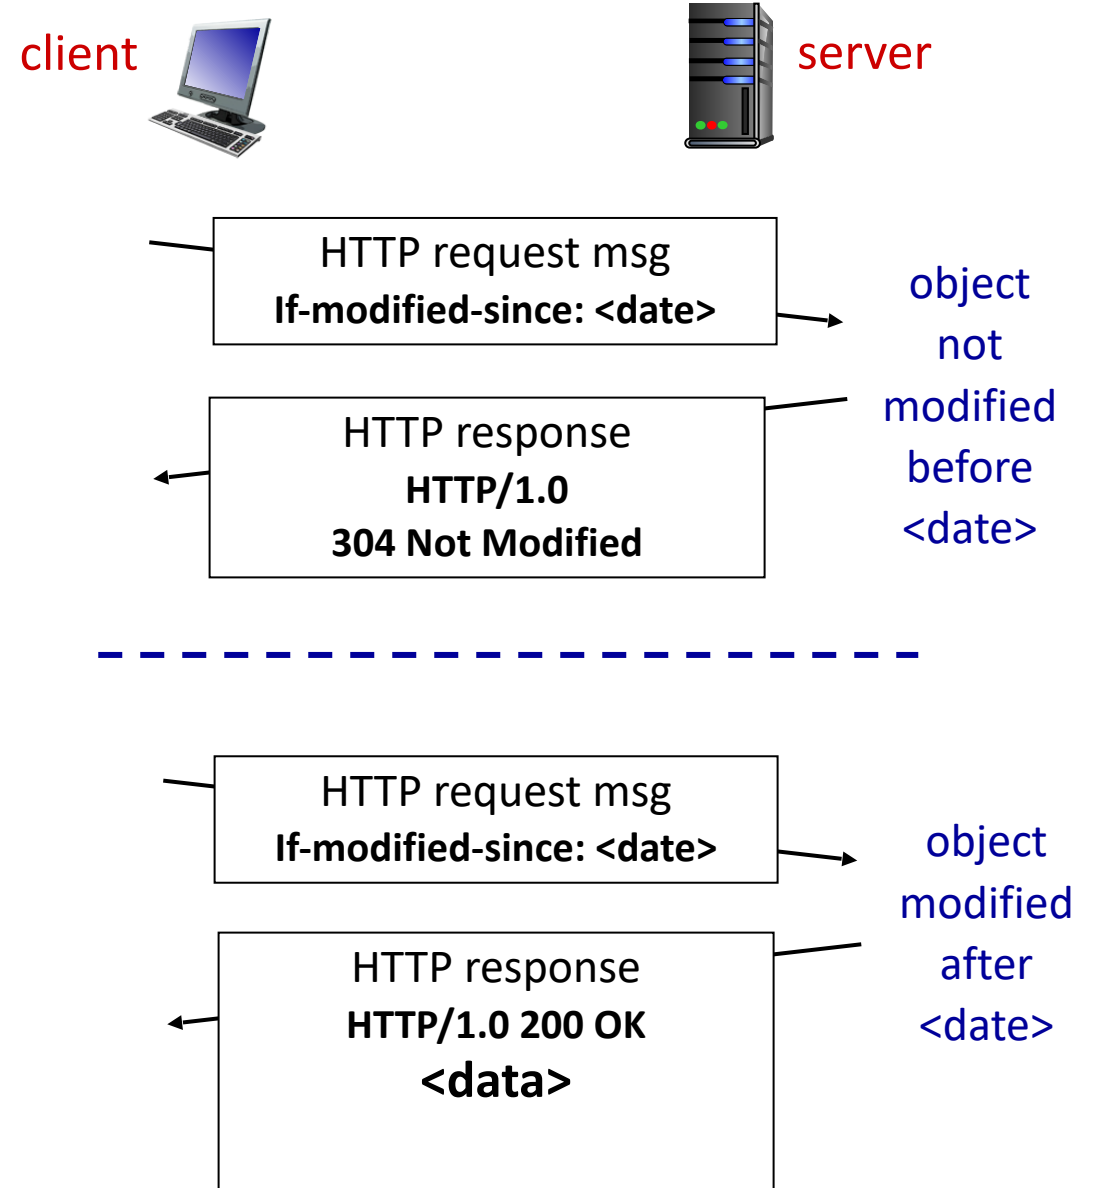
**Goal:** don't send object if cache has up-to-date cached version

- no object transmission delay (or use of network resources)

- **client:** specify date of cached copy in HTTP request

**If-modified-since: <date>**

- **server:** response contains no object if cached copy is up-to-date:  
**HTTP/1.0 304 Not Modified**



# Resumo de Proxy

HTTP (Cookies)

Proxy

- ☐ Atende o cliente sem envolver o servidor Web
- ☐ Transparente × Configurado pelo usuário
- ☐ Utilidade
  - Cache
  - Controle de acesso
  - Acesso a conteúdo restrito para determinados países (IPs)
- ☐ Pelo menos dois sockets envolvidos
  - 1 socket do cliente para o servidor proxy
  - 1 socket do servidor proxy para o servidor HTTP